

RAJALAKSHMI ENGINEERING COLLEGE, CHENNAI

(Autonomous Institution)

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

PHASE I REPORT

ON

**FINANCIAL STATEMENT ANALYTICS
PIPELINE**

Submitted by

PADMASRI K M (230701223)

POOJA A K (230701226)

RAGHUL R (230701251)

BACHELOR OF ENGINEERING

IN

COMPUTER SCIENCE AND ENGINEERING



Rajalakshmi Engineering College, Chennai

Affiliated to Anna University, Chennai

2026

ANNA UNIVERSITY, CHENNAI

BONAFIDE CERTIFICATE

Certified that this Report titled “FINANCIAL STATEMENT ANALYTICS PIPELINE “ is the bonafide work of K M PADMASRI (230701223), A K POOJA (230701226), R RAGHUL (230701251) who carried out the work under my supervision. Certified further that to the best of my knowledge the work reported herein does not form part of any other thesis or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidates.

SIGNATURE

**Dr. Manoranjini,
Professor and Head,
Department of Computer Science
and Engineering,
Rajalakshmi Engineering College,
Chennai – 602 105.**

SIGNATURE

**Mrs. B. Vani ,
Assistant Professor
Department of Computer Science
and Engineering,
Rajalakshmi Engineering College,
Chennai – 602 105.**

Submitted to Project Viva-Voce Examination held on 07.10.2026

Internal Examiner

External Examiner

ABSTRACT

The Financial Statement Analytics Pipeline is an integrated intelligent system developed to automate the extraction, preprocessing, storage, analysis, and interpretation of financial statement data. Conventional financial analysis often requires users to manually collect values from financial reports, clean inconsistent data, calculate financial ratios, compare performance across different periods, and identify potential financial risks. These manual processes can be time-consuming, repetitive, and prone to calculation and interpretation errors. The proposed system aims to overcome these limitations by providing an end-to-end automated framework for financial statement analysis.

The system begins with the upload of financial statements, from which relevant financial information is extracted using document-processing and data-extraction techniques. The extracted data is then subjected to cleaning and standardization, where inconsistent formats, missing values, duplicate entries, and variations in financial identifiers are handled to produce reliable and structured financial data. The processed information is stored in a database for efficient retrieval and further analysis. Based on the standardized financial data, the system automatically calculates key financial ratios and performance indicators, including profitability, liquidity, solvency, and efficiency measures. It also performs financial trend analysis by comparing important metrics across different reporting periods, enabling users to understand changes in the organization's financial performance.

To enhance financial risk assessment, the pipeline incorporates anomaly and risk detection techniques to identify unusual patterns and potentially concerning changes in financial data. The system further provides a What-If Simulator, allowing users to modify selected financial parameters such as revenue, operating expenses, assets, liabilities, or other relevant values and observe their potential impact on financial ratios and overall financial performance. This scenario-based analysis helps users evaluate possible outcomes under different financial conditions without modifying the original stored financial data.

The pipeline also incorporates Explainable AI to convert analytical and risk-detection results into understandable insights. Instead of presenting only numerical outputs, the system provides explanations regarding significant changes, detected anomalies, financial trends, and possible risk factors. Finally,

an interactive financial analytics dashboard presents the extracted data, calculated ratios, trends, detected risks, What-If simulation results, and generated insights through easily understandable visualizations. By integrating automated data processing, financial analysis, risk detection, scenario simulation, and explainable insights into a single pipeline, the proposed system aims to reduce manual effort, improve analytical efficiency and consistency, and support more informed and data-driven financial decision-making.

ACKNOWLEDGEMENT

Initially we thank the Almighty for being with us through every walk of our lives and showering his blessings through the endeavor to put forth this report.

Sincere thanks to our Chairman, Mr. S. MEGANATHAN, B.E., F.I.E., our Vice Chairman, Mr. ABHAY SHANKAR MEGANATHAN, B.E., M.S., and our respected Chairperson, Dr. (Mrs.) THANGAM MEGANATHAN, Ph.D., for providing us with the requisite infrastructure and sincere endeavor in educating us in their premier institution.

Our sincere thanks to Dr. S.N. MURUGESAN, M.E., Ph.D., our beloved Principal for his kind support and facilities provided to complete our work in time.

We express our sincere thanks to Dr. MANORANJINI, Professor and Head of the Department of Computer Science and Engineering for his guidance and encouragement throughout the project work.

We convey our sincere and deepest gratitude to our internal guide, Mrs. B. VANI , Assistant Professor, Department of Computer Science and Engineering, for his valuable guidance throughout the course of the project.

We thank our Project Coordinator, Mr. Madhusudhanan, Assistant Professor, Department of Computer Science and Engineering, for his useful tips during our review to build our project.

K M PADMASRI

A K POOJA

R RAGHUL

TABLE OF CONTENTS

ABSTRACT iii

LIST OF TABLES viii

LIST OF FIGURES ix

LIST OF ABBREVIATIONS x

1. INTRODUCTION 1
 - 1.1 GENERAL 1
 - 1.2 OBJECTIVE 2
 - 1.3 EXISTING SYSTEM 2
 - 1.4 PROPOSED SYSTEM 3
2. LITERATURE SURVEY 6
 - 2.1 LITERATURE REVIEW 5
 - 2.2 COMPARISON AND DISCUSSION 11
 - 2.3 CONCLUSION 12
3. SYSTEM DESIGN 13
 - 3.1 SYSTEM ARCHITECTURE 13
 - 3.2 SYSTEM REQUIREMENTS 16
 - 3.2.1 SOFTWARE REQUIREMENTS 16
 - 3.2.2 HARDWARE REQUIREMENTS 18
 - 3.2.3 DATABASE REQUIREMENTS 19
 - 3.2.4 DEPLOYMENT AND SCALING 20
 - 3.3 SYSTEM DESIGN 20
 - 3.3.1 ACTIVITY DIAGRAM 20
 - 3.3.2 DATA FLOW DIAGRAM 23
4. PROJECT DESCRIPTION 24
 - 4.1 METHODOLOGIES 24
 - 4.2 MODULES 24
 - 4.2.1 DATA MANAGER PACKAGE 25
 - 4.2.2 CORE PACKAGE 26
 - 4.2.3 MODERATION PACKAGE 28
 - 4.2.4 REINFORCEMENT PACKAGE 30
 - 4.2.5 UI PACKAGE 30

5.	IMPLEMENTATION AND RESULT DISCUSSION	31
5.1	IMPLEMENTATION RESULTS	31
5.1.1	IMPLEMENTATION OF DATA PRE-PROCESSING	31
5.1.2	IMPLEMENTATION OF DATA TOKENIZATION	36
5.1.3	IMPLEMENTATION OF DATA EMBEDDING	37
5.2	DISCUSSION	38
6.	CONCLUSION AND FUTURE WORK	40
6.1	CONCLUSION	40
6.2	FUTURE WORK	40
	REFERENCES	42

INTRODUCTION

1.1 GENERAL

Financial statements are an essential source of information for understanding the financial condition and performance of an organization. Statements such as the **Balance Sheet, Income Statement, Cash Flow Statement, and related financial reports** contain important information about an organization's assets, liabilities, equity, revenue, expenses, cash flows, profitability, and overall financial position. These statements are used by management, investors, financial analysts, creditors, auditors, and other stakeholders to evaluate the organization's performance and make informed decisions.

However, analyzing financial statements manually can be a time-consuming and repetitive process, particularly when information is available in large PDF reports or documents with different structures and formats. Financial values may be represented using different units, naming conventions, date formats, and layouts. Extracting the required information manually and entering it into spreadsheets can introduce errors and make the analysis difficult to maintain. In addition, calculating multiple financial ratios, comparing performance across several periods, identifying unusual changes, and interpreting the results requires considerable effort and financial knowledge.

The **Financial Statement Analytics Pipeline** is proposed as an integrated solution to automate these activities. The system provides an end-to-end workflow beginning with the **upload of financial statements** and continuing through financial-data extraction, preprocessing, standardization, storage, analysis, risk detection, scenario simulation, and visualization. The pipeline is designed to transform unstructured or semi-structured financial documents into structured and meaningful financial information that can be used for further analysis.

The first stage of the pipeline focuses on **financial data extraction**. Relevant values and financial identifiers are obtained from uploaded financial statements using document-processing techniques. Since financial documents may contain inconsistent formats, the extracted information is passed through a **data cleaning and standardization stage**. This stage handles issues such as inconsistent labels, numerical formats, missing values, duplicate records, and

differences in units. The objective is to produce reliable and standardized financial data before performing calculations or analytical operations.

After preprocessing, the cleaned financial data is stored in a **structured database**. Persistent storage allows the system to maintain historical financial information and retrieve it whenever required for analysis. The stored data forms the foundation for calculating financial indicators and comparing organizational performance across different reporting periods.

The system then performs **financial ratio and trend analysis**. Important categories of financial ratios, including **liquidity, profitability, solvency, and efficiency ratios**, can be calculated from the stored financial data. Trend analysis allows users to compare financial performance over multiple periods and identify increases, decreases, or significant changes in important financial indicators. Instead of requiring users to perform repeated calculations manually, the pipeline automatically generates these analytical results.

Another important component of the system is **financial risk and anomaly detection**. Financial data can contain unusual changes or patterns that may require further investigation. The proposed pipeline applies analytical and machine-learning techniques to identify potential anomalies and risk indicators. This provides an additional layer of analysis beyond conventional ratio calculations and helps users focus their attention on potentially significant financial changes.

The pipeline also incorporates a **What-If Simulator** for scenario-based financial analysis. Users can modify selected financial parameters, such as revenue, expenses, assets, liabilities, or other relevant values, and observe how these changes may affect financial ratios and overall financial performance. The simulator allows different hypothetical scenarios to be evaluated without changing the original financial records. This can help users understand the potential consequences of financial decisions and compare alternative situations before making decisions.

To make the analytical results easier to understand, the system incorporates **Explainable AI and financial insights**. Complex numerical results and detected patterns can be converted into meaningful explanations that help users understand why a particular financial indicator has changed, why an anomaly has been identified, or which factors may contribute to a particular risk. This

improves the interpretability of the system and reduces the dependence on users having advanced technical or analytical knowledge.

Finally, the results are presented through an **interactive financial analytics dashboard**. The dashboard provides a centralized view of financial metrics, ratio calculations, historical trends, detected anomalies, risk indicators, What-If simulation results, and generated insights. Visual representations such as charts, graphs, tables, and key performance indicators make it easier to interpret large amounts of financial information and support data-driven decision-making.

Overall, the **Financial Statement Analytics Pipeline** aims to provide a complete and systematic approach to financial statement analysis by integrating document processing, data preprocessing, database management, financial calculations, trend analysis, anomaly detection, scenario simulation, explainable AI, and visualization into a single system. The proposed approach reduces repetitive manual work, improves consistency in financial analysis, enables faster interpretation of financial information, and provides users with a structured platform for evaluating financial performance and potential risks.

1.2 OBJECTIVE OF THE PROJECT

The project aims to develop an integrated **Financial Statement Analytics Pipeline** that automates financial data extraction, cleaning, standardization, and storage while reducing the effort involved in manual analysis. It provides automated **financial ratio and trend analysis** to evaluate organizational performance and uses analytical techniques to identify potential risks and unusual financial patterns. A **What-If Simulator** enables users to modify key financial parameters and understand their possible impact on financial performance and ratios. The system further provides explainable financial insights and an interactive dashboard, supporting efficient, accurate, and data-driven financial decision-making.

1.3 EXISTING SYSTEM

The existing approach to financial statement analysis mainly depends on **manual data extraction, spreadsheet-based calculations, and separate financial analysis tools**. Financial information is generally collected from annual reports and financial statements and then manually entered into

spreadsheets for further processing. Financial ratios such as profitability, liquidity, and solvency ratios are calculated separately, while comparisons between different financial periods are often performed manually.

Although spreadsheet applications can support basic calculations and visualization, they provide limited automation for **financial data extraction, anomaly detection, risk analysis, and intelligent interpretation**. Financial statements may also contain different formats, layouts, units, and terminology, making manual data preparation time-consuming and prone to errors. Identifying unusual financial patterns generally requires additional analysis and financial expertise.

Furthermore, conventional systems usually do not provide an integrated **What-If Simulator** for evaluating hypothetical changes in financial parameters. Users may need to manually modify values and recalculate ratios to understand possible outcomes. Explainable AI-based insights are also limited, making it difficult for users to understand the reasons behind detected financial risks or significant changes. Therefore, the existing system lacks a unified platform that combines **financial data extraction, preprocessing, storage, ratio and trend analysis, risk detection, What-If simulation, explainable insights, and interactive visualization** in a single automated pipeline.

1.4 PROPOSED SYSTEM

The proposed **Financial Statement Analytics Pipeline** provides an integrated and automated approach for processing and analyzing financial statements. The system begins with the **upload of financial statement documents**, followed by automated extraction of relevant financial information using document-processing and OCR techniques. The extracted data is then cleaned, validated, and standardized to handle missing values, inconsistent formats, and variations in financial identifiers before being stored in a structured database.

The system automatically performs **financial ratio and trend analysis** to evaluate profitability, liquidity, solvency, and overall financial performance across different reporting periods. Machine-learning techniques are incorporated for **financial risk and anomaly detection**, helping identify unusual patterns that may require further investigation. A **What-If Scenario Simulator** allows users to modify important financial parameters and observe their potential

impact on financial ratios and overall performance without affecting the original financial data.

The proposed system also incorporates **Explainable AI** to provide understandable interpretations of analytical results, detected anomalies, and financial risks. An **interactive dashboard** presents financial metrics, trends, risk indicators, simulation results, and generated insights through charts, graphs, and tables. By combining data extraction, preprocessing, storage, financial analysis, risk detection, What-If simulation, explainable insights, and visualization into a single pipeline, the proposed system reduces manual effort, improves analytical efficiency, and supports faster and more informed financial decision-making.




Figure 1.1 – System Architecture of the Financial Statement Analytics Pipeline

As per the diagram, the **Financial Statement Analytics Pipeline** follows an end-to-end architecture that processes financial information from multiple sources and converts it into meaningful financial insights. The system accepts data from **PDF reports, XBRL filings, Excel/CSV files, web sources, databases, APIs, and other third-party sources**. The collected information first passes through the **Ingestion and Extraction Layer**, where documents are processed, text and tables are extracted, OCR and table detection are performed, and relevant financial fields are identified and classified. The extracted information is then transferred to the **Data Processing and Standardization Layer**, where data cleaning, validation, standardization, entity mapping, and

quality checks are performed to ensure consistency and reliability. The processed financial data is maintained in the **Data Storage Layer**, which contains financial fact tables, company master information, statement metadata, audit logs, versioning, and data lineage. The stored data is subsequently analyzed through the **Analytics and Intelligence Engine**, which performs financial calculations, feature engineering, machine-learning-based scoring, rule-based analysis, and explainability. The system provides major analytical features such as **financial ratio and trend analysis, financial health scoring, anomaly and fraud risk detection, profit–cash flow and cross-statement consistency analysis, financial distress forecasting, explainable AI and evidence-based insights, and What-If financial scenario simulation**. The generated results are presented through the **Output and Access Layer** using dashboards, visualizations, reports, insights, alerts, notifications, the scenario simulator interface, and API access. The entire architecture is supported by an **Infrastructure and Security Layer** that provides secure storage, access control, encryption, monitoring, logging, backup, and recovery. The final analytical results can be utilized by **investors, analysts, auditors, and management** to understand financial performance, assess risks, evaluate possible scenarios, and make informed financial decisions.

CHAPTER 2

LITERATURE SURVEY

2.1 LITERATURE SURVEY

Financial statement analysis has traditionally relied on ratio analysis, comparative analysis, forecasting, and manual interpretation of accounting information. Recent research has increasingly focused on automating financial-data extraction, standardization, fraud detection, risk prediction, natural language processing, and intelligent financial analysis. The following studies provide the foundation for the development of the proposed **Financial Statement Analytics Pipeline**.

2.1.1 Financial Big Data Analysis and Early Warning Platform: A Case Study – Liang et al. (2020)

Liang et al. presented a big-data-based platform for financial risk analysis and early warning. The study combines heterogeneous data collection, data cleaning, data fusion, cloud computing, and artificial intelligence to improve financial risk monitoring and decision-making. The work demonstrates the importance of integrating different financial data sources and applying automated processing techniques for efficient financial analysis. It also highlights the challenges associated with data quality, integration, confidentiality, and the complexity of financial-data analysis. The study provides a foundation for incorporating automated data processing and risk-analysis capabilities into the proposed financial analytics pipeline.

2.1.2 Intelligent Fraud Detection in Financial Statements Using Machine Learning and Data Mining – Ashtiani and Raahemi (2021)

Ashtiani and Raahemi conducted a systematic review of machine-learning and data-mining techniques used for detecting fraudulent financial statements. The study analyzed 47 research papers and examined different approaches for identifying financial statement fraud. The authors highlighted the importance of anomaly detection and the use of unstructured financial information along with advanced machine-learning methods. The findings demonstrate that intelligent techniques can improve the identification of suspicious financial patterns and provide a basis for incorporating machine-learning-based risk and anomaly detection into financial statement analytics.

2.1.3 Financial Knowledge Graph Based Financial Report Query System – Zehra et al. (2021)

Zehra et al. developed a financial knowledge graph-based system for extracting and querying information from annual financial reports. The approach uses ontology-based information extraction to identify useful financial information and represent relationships between financial entities. The system enables users to query financial information and automatically generate financial stories that can support decision-making. This study demonstrates the usefulness of automated information extraction and structured representation of financial-report content, which is relevant to the extraction and standardization stages of the proposed Financial Statement Analytics Pipeline.

2.1.4 Sentiment Analysis in Finance: From Transformers Back to eXplainable Lexicons (XLex) – Rizinski et al. (2024)

Rizinski et al. proposed XLex, an explainable financial sentiment-analysis approach that combines transformer-based models with SHAP-based explainability. The study focuses on improving financial vocabulary coverage and sentiment classification while maintaining interpretability and computational efficiency. The work demonstrates how domain-specific NLP and explainable AI can be applied to financial information to generate more meaningful insights. Its approach is relevant to the proposed system's Explainable AI component, particularly for providing understandable explanations rather than presenting only numerical analytical results.

2.1.5 AI Driven Fraud Detection Models in Financial Networks: A Comprehensive Systematic Review – Sarna et al. (2025)

Sarna et al. presented a systematic review of AI-driven fraud detection models used in financial networks. The study examines machine learning, deep learning, hybrid models, and graph neural networks for identifying different types of financial fraud. The review demonstrates the ability of AI techniques to identify complex fraud patterns while also highlighting important challenges such as model explainability, privacy, bias, and scalability. These findings support the inclusion of intelligent risk and anomaly detection in the proposed pipeline while emphasizing the importance of explainable and reliable financial analytics.

2.1.6 Multimodal Fraud Detection in Financial Statements: A Trimodal Attention Network With Contrastive Evidence Chain Construction – Ping et al. (2026)

Ping et al. proposed **TriAudit**, a multimodal framework for financial statement fraud detection. The approach combines structured financial ratios, textual information, and document-layout features to identify fraudulent patterns. By integrating different types of financial evidence, the system improves fraud detection while providing interpretable evidence linking financial, textual, and visual anomalies. The study demonstrates the advantages of multimodal financial analysis and supports the proposed pipeline's goal of combining structured financial information with advanced analytical and explainable techniques.

2.1.7 Mining Corporate Annual Reports for Intelligent Detection of Financial Statement Fraud – Hajek and Henriques (2017)

Hajek and Henriques investigated the use of machine-learning techniques for detecting financial statement fraud from corporate annual reports. The study combines structured financial information with managerial comments and applies machine-learning methods to identify fraudulent patterns. The integration of numerical financial data and textual information is particularly relevant to automated financial analysis because annual reports contain both quantitative and qualitative information. The study supports the proposed pipeline's approach of combining structured financial data with unstructured textual information for more comprehensive financial analysis.

2.1.8 A Research on Text Information Extraction from Annual Report Based on Domain Ontology (2017)

This study applied domain ontology and natural-language-processing techniques to extract useful information from annual financial reports. The approach focuses on understanding financial-domain terminology and extracting relevant information from textual reports. Such automated extraction can reduce the need for manual reading and data entry when processing large financial documents. The study is directly relevant to the proposed pipeline's financial-document extraction and standardization stages, where extracted information must be mapped into a consistent structure before further analysis.

2.1.9 Segment Information Extraction from Financial Annual Reports Using Neural Network – Ito, Sakaji and Izumi (2019)

Ito, Sakaji, and Izumi investigated the use of neural-network techniques for extracting segment information from financial annual reports. The study demonstrates how machine-learning models can automatically identify and extract structured financial information from otherwise unstructured annual reports. This approach is useful for reducing manual data extraction and improving the availability of financial information for subsequent analysis. The research supports the automated extraction component of the proposed Financial Statement Analytics Pipeline.

2.1.10 FinBERT: Financial Sentiment Analysis with Pre-trained Language Models – Araci (2019)

Araci introduced **FinBERT**, a BERT-based language model specialized for financial-domain sentiment analysis. The model is designed to understand financial language more effectively than general-purpose language models by using financial-domain knowledge. The study provides a foundation for applying domain-specific NLP techniques to financial documents and extracting meaningful information from financial text. This work is relevant to the proposed pipeline's Explainable AI and financial-insight generation components, particularly where textual financial information needs to be interpreted.

2.1.11 FinQA: A Dataset of Numerical Reasoning over Financial Data – Chen et al. (2021)

Chen et al. introduced **FinQA**, a dataset designed for numerical reasoning over financial reports. The dataset requires systems to combine information from financial reports with numerical reasoning to answer financial questions. This work demonstrates the importance of reasoning over financial data rather than simply extracting individual values. It is relevant to the proposed system because financial analytics requires relationships between extracted values, financial calculations, and meaningful interpretations. The study therefore provides support for the numerical reasoning and AI-based financial insight components of the proposed pipeline.

2.1.12 TAT-QA: A Question Answering Benchmark on a Hybrid of Tabular and Textual Content in Finance – Zhu et al. (2021)

Zhu et al. developed **TAT-QA**, a financial question-answering benchmark that combines tabular and textual information. The work requires numerical reasoning over both financial tables and narrative content, demonstrating that effective financial analysis may require information from multiple representations within the same report. This study supports the integration of extracted financial tables and textual information in the proposed analytics pipeline. It also highlights the need for systems capable of processing and reasoning over both structured and unstructured financial information.

2.1.13 On the Determinants and Prediction of Corporate Financial Distress in India (2021)

This study investigated the prediction of corporate financial distress in the Indian context using financial ratios and machine-learning models such as **Support Vector Machines (SVM)** and **Artificial Neural Networks (ANN)**. The research demonstrates that financial indicators can be used as meaningful inputs for predicting potential financial distress. The study is particularly relevant to the proposed Financial Statement Analytics Pipeline because it provides evidence for combining conventional financial indicators with machine-learning techniques for financial risk prediction.

2.1.14 Predicting Financial Distress of Chinese Listed Companies Using Machine Learning: To What Extent Does Textual Disclosure Matter? (2023)

This study examined whether textual disclosures, when combined with conventional financial information, can improve financial distress prediction. The research demonstrates the value of integrating structured financial metrics with unstructured textual information when evaluating the financial condition of companies. This supports the proposed pipeline's approach of extending beyond traditional numerical ratio analysis and incorporating additional information sources for risk assessment. The findings are particularly relevant to the proposed financial risk detection and Explainable AI components.

2.1.15 Assessing Large Language Models Used for Extracting Table Information from Annual Financial Reports (2024)

This study investigated the application of **Large Language Models (LLMs)** for extracting and interpreting tabular information from annual financial reports. The research highlights the potential of modern AI techniques for automating

financial-document and table extraction, which is traditionally a time-consuming manual process. The study is directly relevant to the proposed Financial Statement Analytics Pipeline because automated extraction is the initial stage of the system. The use of modern AI-based extraction techniques can help transform financial reports into structured data that can subsequently undergo cleaning, standardization, storage, ratio analysis, risk detection, and financial insight generation.

2.2 COMPARISON AND DISCUSSION

The reviewed literature shows that existing financial analytics research focuses on different aspects of financial statement processing and analysis. Liang et al. focused on **financial big-data analysis and early warning**, demonstrating the importance of data integration, cleaning, and AI-based risk monitoring. Ashtiani and Raahemi concentrated on **machine-learning and data-mining techniques for financial statement fraud detection**, while Sarna et al. further explored AI, deep learning, hybrid models, and GNN-based approaches for financial fraud detection.

A different group of studies concentrated on **financial-document information extraction**. Zehra et al. used ontology-based information extraction and financial knowledge graphs for querying annual reports, while the 2017 domain-ontology study and Ito et al. explored automated extraction of information from annual reports using NLP and neural networks. More recent work on large language models also investigates automated extraction and interpretation of tabular information from financial reports. These studies demonstrate the importance of converting unstructured financial documents into structured information before performing analysis.

Several studies also demonstrate the importance of **combining structured and unstructured financial information**. Chen et al. introduced FinQA for numerical reasoning over financial reports, while Zhu et al. developed TAT-QA to combine tabular and textual financial information. Similarly, research on financial distress prediction shows that combining financial indicators with textual disclosures can improve the assessment of financial conditions. Ping et al. extended this concept further by combining financial ratios, textual information, and document-layout features for multimodal fraud detection.

The literature also indicates a growing requirement for **explainable financial AI**. Rizinski et al. incorporated SHAP explainability into financial sentiment

analysis, while other fraud-detection studies identify explainability as an important challenge for AI-based financial systems. FinBERT further demonstrates the usefulness of domain-specific language models for understanding financial text. These works indicate that financial analytics systems should not only produce predictions or risk scores but should also provide understandable explanations for their results.

In comparison, the **proposed Financial Statement Analytics Pipeline** aims to integrate the major capabilities identified across the reviewed studies into a single workflow. Rather than focusing only on fraud detection, document extraction, sentiment analysis, or financial distress prediction, the proposed system connects **financial-data extraction, cleaning, standardization, storage, financial ratio and trend analysis, anomaly and risk detection, What-If scenario simulation, Explainable AI, and interactive dashboard visualization**. The project therefore focuses on the integration of multiple analytical stages rather than introducing a completely new financial ratio or a single isolated prediction model. This integrated approach is consistent with the source paper's emphasis on connecting data preparation with financial analysis and reducing repetitive manual operations.

Comparison of Reviewed Studies

Study	Main Focus	Major Technique	Limitation / Gap	Relevance to Proposed System
Liang et al. (2020)	Financial risk and early warning	Big data, AI, data fusion	Mainly focuses on risk monitoring	Supports data integration and risk analysis
Ashtiani & Raahemi (2021)	Financial statement fraud	ML and data mining	Focuses mainly on fraud detection	Supports anomaly and fraud detection
Zehra et al. (2021)	Financial report querying	Knowledge graph, ontology	Focuses on report querying	Supports financial information extraction

Rizinski et al. (2024)	Financial sentiment analysis	Transformers, SHAP	Focuses on textual sentiment	Supports Explainable AI
Sarna et al. (2025)	AI-based financial fraud	ML, DL, GNNs	Highlights explainability and scalability challenges	Supports intelligent risk detection
Ping et al. (2026)	Multimodal fraud detection	Ratios, text, layout, attention	Primarily focused on fraud	Supports multimodal financial analysis
Hajek & Henriques (2017)	Fraud detection from annual reports	ML + financial/text data	Focused on fraud detection	Supports structured + textual analysis
Ito et al. (2019)	Financial information extraction	Neural networks	Focused on segment extraction	Supports automated extraction
Araci (2019)	Financial NLP	FinBERT	Mainly sentiment analysis	Supports financial-domain NLP
Chen et al. (2021)	Financial numerical reasoning	FinQA	Focuses on question answering	Supports financial reasoning
Zhu et al. (2021)	Tabular + textual finance QA	TAT-QA	Focuses on QA	Supports combined financial data analysis
Financial Distress Study (2021)	Financial distress prediction	SVM, ANN, financial ratios	Focused on distress prediction	Supports risk prediction

Chinese Financial Distress Study (2023)	Financial distress + text	ML + textual disclosures	Specific prediction task	Supports structured + unstructured analysis
LLM Table Extraction Study (2024)	Financial table extraction	LLMs	Focuses on extraction	Supports automated document processing

Table 1.1 Comparison of literature survey papers

Overall, the comparison indicates that **no single reviewed study covers the complete pipeline proposed in this project**. The proposed system brings together the major concepts identified in the literature and extends them with a **What-If Simulator**, enabling users to evaluate hypothetical changes in financial variables and observe their effect on financial metrics. This makes the system more suitable as an integrated financial analysis and decision-support platform.

2.3 CONCLUSION

The literature survey demonstrates the increasing use of automation, machine learning, natural language processing, and artificial intelligence in financial statement analysis. The reviewed studies have addressed important areas such as financial-data extraction, fraud detection, financial risk assessment, financial distress prediction, numerical reasoning, textual analysis, and explainable AI. These studies show that automated techniques can reduce manual effort and improve the efficiency of processing and analyzing financial information.

However, most existing approaches concentrate on a specific financial-analysis task rather than providing a complete end-to-end analytical workflow. The reviewed research also highlights challenges related to data quality, inconsistent financial formats, unstructured documents, explainability, and the integration of structured and textual information. The proposed Financial Statement Analytics Pipeline addresses these gaps by integrating financial-data extraction, cleaning, standardization, storage, ratio and trend analysis, risk and anomaly detection,

What-If scenario simulation, Explainable AI, and interactive visualization into a unified system.

Thus, the literature provides a strong foundation for developing an integrated financial analytics solution that can transform raw financial information into meaningful and interpretable insights. The proposed system extends the existing approaches by combining multiple stages of financial analysis and providing scenario-based evaluation through the What-If Simulator, thereby supporting more efficient and informed financial decision-making.

CHAPTER 3

SYSTEM DESIGN

3.1 SYSTEM DESIGN

The **Financial Statement Analytics Pipeline** is designed as an end-to-end intelligent financial analytics system that converts financial information from multiple sources into structured, reliable, and actionable insights. The system follows a layered architecture consisting of **data sources, ingestion and extraction, data processing and standardization, data storage, analytics and intelligence, output and access, and infrastructure and security**. The architecture supports financial statement extraction, financial analytics, risk assessment, forecasting, explainable AI, and What-If scenario simulation.

3.1.1 Data Sources

The pipeline accepts financial information from different sources, including **PDF reports, annual reports, XBRL filings, Excel/CSV files, web sources, databases, APIs, and third-party sources**. Supporting multiple sources enables the system to process financial information available in different formats and structures.

3.1.2 Ingestion and Extraction Layer

The ingestion layer receives the financial documents and performs automated extraction. **Python** is used as the primary programming language, while **OpenCV and Tesseract OCR** support OCR-based extraction. **PyMuPDF and pdfplumber** are used for PDF processing and table extraction, while **Camelot and Tabula** can be used for extracting tabular financial information. For XBRL-based financial filings, **XBRL parsing libraries such as Arelle** are included. The extracted information is mapped into relevant financial fields and classified before producing the raw data output.

3.1.3 Data Processing and Standardization Layer

The extracted information is processed using **Pandas and NumPy**. This stage performs data cleaning, validation, standardization, entity mapping, and quality checking. Missing values, inconsistent formats, duplicate information, and variations in financial terminology are handled before the data is used for

analysis. Standardization ensures that financial information from different companies and reporting periods can be compared consistently.

3.1.4 Data Storage Layer

The standardized financial information is stored in a relational database using **PostgreSQL or MySQL**. **SQLAlchemy** is used for database connectivity. The storage layer maintains financial facts, company information, statement metadata, and other processed information required for subsequent analysis. **Audit logs, versioning, and data lineage** can also be maintained to improve traceability and reliability.

3.1.5 Analytics and Intelligence Engine

The analytics and intelligence engine performs the core financial analysis. **Pandas and NumPy** support financial calculations and data processing, while **Scikit-learn** is used for machine-learning models. **XGBoost** is included for advanced risk and anomaly-related analysis. The system performs financial ratio calculation, trend analysis, financial health scoring, anomaly and fraud-risk detection, cross-statement consistency analysis, and financial distress forecasting. The architecture also includes a rule engine and explainability engine for interpreting analytical results.

3.1.6 Financial Health and Risk Analysis

The system generates a **Financial Health Score ranging from 0–100** using important financial indicators. It also identifies unusual patterns and potential fraud risks through anomaly detection. Cross-statement analysis can identify inconsistencies between the **Profit and Loss Statement, Balance Sheet, and Cash Flow Statement**, while machine-learning models are used for financial distress forecasting and early-warning analysis.

3.1.7 What-If Financial Scenario Simulation

The system provides a **What-If Financial Scenario Simulator** that allows users to modify financial assumptions such as **revenue, expenses, and debt** and observe their possible impact on financial metrics and the overall financial health score. A scenario engine and ML models can be used to support this functionality. This feature enables users to compare hypothetical situations and understand the potential consequences of financial decisions.

3.1.8 Explainable AI

The Explainable AI component provides evidence-based explanations for predictions, risk scores, anomalies, and financial changes. **SHAP and LIME** are identified as the explainability technologies in the project. NLP techniques can additionally support the interpretation of financial text. This component improves transparency by helping users understand the factors responsible for analytical results rather than presenting only a prediction or score.

3.1.9 Output and Visualization Layer

The processed results are presented through **interactive dashboards, reports, insights, alerts, notifications, scenario simulation interfaces, and API access**. **Streamlit, Plotly, and Power BI** are identified for visualization, while the project architecture also lists **React/Streamlit and Plotly** for the interactive dashboard. These tools allow users to view financial ratios, trends, health scores, risk indicators, forecasts, and What-If simulation results in an understandable format.

3.1.10 Pipeline Orchestration and Development Tools

Apache Airflow is included as an optional technology for pipeline orchestration, allowing different processing stages to be scheduled and coordinated. **Git and GitHub** are used for version control, while **VS Code and PyCharm** are listed as development environments. **Docker** is included as an optional technology for containerization and deployment.

3.1.11 Datasets

The system can use **sample annual reports in PDF format, XBRL financial filings such as MCA India and SEC EDGAR data, public financial datasets from Kaggle, and company financial statements including Balance Sheets, Income Statements, and Cash Flow Statements**. These datasets provide the financial information required for extraction, preprocessing, analysis, model development, and evaluation.

3.1.12 Overall System Flow

The complete system follows the sequence:

Financial Data Sources → Data Ingestion → Text/Table Extraction → OCR & Table Detection → Field Mapping & Classification → Data Cleaning → Validation → Standardization → Entity Mapping → Quality Checks → Database Storage → Financial Calculations → ML Models &

Scoring → Rule Engine → Explainability → What-If Simulation → Dashboard, Reports, Alerts and Insights.


This architecture provides a modular structure in which each stage can be developed and improved independently while maintaining the complete end-to-end financial analytics workflow. The methodology presented in the PPT specifically organizes the project into **data extraction, cleaning and validation, standardization and storage, financial analytics, ML/AI analysis, What-If simulation, Explainable AI, and visualization and insights.**

Technologies Used

Category	Technologies
Programming	Python
Data Processing	Pandas, NumPy
PDF Extraction	PyMuPDF, pdfplumber, Camelot, Tabula
OCR	OpenCV, Tesseract OCR
XBRL	Arelle / XBRL Parsing Libraries
Machine Learning	Scikit-learn, XGBoost
Explainable AI	SHAP, LIME
NLP	NLP techniques / FinBERT-related approaches
Database	PostgreSQL / MySQL
Database Connectivity	SQLAlchemy
Visualization	Streamlit, Plotly, Power BI
Dashboard	React / Streamlit
Pipeline Orchestration	Apache Airflow (optional)
Version Control	Git, GitHub
Containerization	Docker (optional)

IDE	VS Code, PyCharm
Data Sources	PDF reports, XBRL, Excel/CSV, web sources, APIs, Kaggle datasets, MCA India, SEC EDGAR

Table 3.1 Technologies used



Dia 3.1 System architecture

3.2 SYSTEM REQUIREMENTS

The Financial Statement Analytics Pipeline requires a system capable of handling PDF/XBRL document processing, OCR, data preprocessing, machine-learning analysis, database operations, and interactive visualization. The PPT specifies the major technologies and tools, but it **does not specify exact hardware configurations or software versions**. Therefore, the versions below are recommended compatible versions for implementation, while the technology choices are based directly on the project PPT.

3.2.1 SOFTWARE REQUIREMENTS

The software stack is designed around **Python, Pandas, NumPy, OCR and PDF-extraction tools, machine-learning libraries, explainability frameworks, relational databases, visualization platforms, and development/version-control tools.**

S. No.	Software / Technology	Recommended Version	Purpose
1	Python	3.11.x	Primary programming language and backend processing
2	Pandas	2.2.x	Data cleaning, transformation and financial-data processing
3	NumPy	1.26.x	Numerical calculations and data manipulation
4	PyMuPDF	1.24.x	PDF text and document extraction
5	pdfplumber	0.11.x	PDF text and table extraction
6	Camelot	0.11.x	Extraction of tables from PDF documents
7	Tabula	2.x	PDF table extraction
8	OpenCV	4.10.x	Image processing for OCR preprocessing
9	Tesseract OCR	5.x	OCR-based extraction from scanned financial documents
10	Arelle	2.37.x	XBRL parsing and processing
11	Scikit-learn	1.5.x	Machine-learning models and financial-risk analysis
12	XGBoost	2.1.x	Advanced classification, risk and prediction models
13	SHAP	0.46.x	Explainable AI and feature-importance analysis

14	LIME	0.2.x	Explainable machine-learning predictions
15	PostgreSQL	16.x	Relational database for financial data storage
16	MySQL	8.0.x	Alternative relational database
17	SQLAlchemy	2.0.x	Database connectivity and ORM
18	Plotly	5.x	Interactive financial charts and visualizations
19	Streamlit	1.38.x	Interactive analytics dashboard
20	Power BI	Desktop 2024+	Business intelligence and financial reporting
21	React	18.x	Interactive frontend/dashboard development
22	FastAPI	0.11x	API/backend service for system integration
23	Apache Airflow	2.10.x	Optional pipeline orchestration and scheduling
24	Git	2.4x+	Version control
25	GitHub	Cloud	Source-code repository and collaboration
26	Docker	27.x+	Optional containerization and deployment
27	VS Code	Latest stable	Development environment
28	PyCharm	2024.x+	Alternative Python development environment

Table 3.2.1 Software technologies used

The software requirements consist of the programming languages, libraries, databases, machine-learning frameworks, visualization tools, and development utilities required to implement the proposed pipeline. The project primarily uses

Python, Pandas, NumPy, OCR and PDF extraction libraries, Scikit-learn, XGBoost, SHAP, LIME, PostgreSQL/MySQL, Streamlit, Plotly, Git, and Docker. These technologies support the complete workflow from financial-data extraction to analysis and visualization.

3.2.2 . HARDWARE REQUIREMENT

S. No.	Component	Minimum Requirement	Recommended Requirement
1	Processor	Intel Core i5 / AMD Ryzen 5	Intel Core i7 / AMD Ryzen 7
2	RAM	8 GB	16 GB or above
3	Storage	256 GB SSD	512 GB SSD or above
4	GPU	Integrated GPU	NVIDIA GPU, optional for ML workloads
5	Display	1366 × 768	Full HD 1920 × 1080
6	Internet	Stable broadband connection	High-speed broadband
7	Operating System	Windows 10/11, Linux	Windows 11 / Ubuntu Linux
8	Database Storage	Minimum 10 GB free space	50 GB or more depending on financial documents
9	Backup Storage	External/cloud storage	Cloud or external backup

Table 3.2.2 Hardware technologies used

3.2.3 DATASET REQUIREMENTS

S. No.	Dataset / Source	Format	Purpose
1	Company Annual Reports	PDF	Financial statement extraction

2	Financial Statements	PDF / Excel	Balance Sheet, Income Statement and Cash Flow analysis
3	XBRL Financial Filings	XBRL	Structured financial-data extraction
4	MCA India Filings	XBRL / Financial data	Indian company financial analysis
5	SEC EDGAR	XBRL / Financial filings	Financial-data analysis
6	Public Financial Datasets	CSV / Excel	ML model development and testing
7	Kaggle Financial Datasets	CSV / Excel	Training and evaluation
8	Company Financial Statements	PDF / CSV / Excel	Ratio, trend and risk analysis


Table 3.2.4 Dataset requirements

The project uses **sample annual reports, XBRL financial filings, public financial datasets, and company financial statements**, including Balance Sheets, Income Statements, and Cash Flow Statements. These datasets provide the required financial information for extraction, preprocessing, model development, testing, and financial analysis.

3.3.1. ACTIVITY DIAGRAM

The **active environment** represents the user-driven workflow of the Financial Statement Analytics Pipeline, where the analysis is initiated through direct user interaction. The process begins when the user uploads a financial document such as a **PDF, XBRL filing, or Excel file**, or provides financial data through the application. The system extracts and preprocesses the required information using **PyMuPDF, pdfplumber, OCR, and Arelle**, followed by data validation and standardization using **Pandas and NumPy**. The processed data is then passed to the analytics engine for **financial ratio calculation, trend analysis,**


risk detection, and machine-learning-based analysis. The system checks the processed results for potential anomalies or financial risks. When a risk or anomaly is detected, the system generates appropriate alerts and explanations using **SHAP/LIME**; otherwise, the normal financial analysis results such as ratios, trends, and the financial health score are generated. The results are displayed through the interactive dashboard, reports, and insights using **Streamlit, Plotly, or Power BI**. Users can further perform **What-If scenario simulations** and examine explanations before making financial decisions. This environment provides an interactive and immediate workflow in which the user can upload data, analyze results, explore scenarios, and obtain actionable financial insights.



Dia 3.3.1(a) Workflow in active environment

The **passive environment** represents the batch-based or scheduled workflow of the Financial Statement Analytics Pipeline, where processing occurs


automatically without continuous direct user interaction. The process begins when a scheduler such as **Apache Airflow** triggers the pipeline at a predefined interval. Financial information is collected from sources such as **APIs, company websites, databases, and XBRL filings** and temporarily stored for processing. The system performs batch data cleaning, validation, standardization, and entity mapping using **Pandas and NumPy**, after which the processed information is stored in **PostgreSQL or MySQL**. The analytics engine then performs **financial ratio and trend analysis, financial health scoring, anomaly and risk detection, and financial distress forecasting** using machine-learning techniques. The system checks for newly available financial data and periodically updates the analytical results when new information is received. Updated dashboards, reports, and insights are generated using **Plotly, Streamlit, or Power BI**, while notifications and alerts can be issued when required. This passive environment enables the pipeline to continuously maintain updated financial information and analytical results through scheduled automated processing.



Dia 3.3.1 (b) Workflow in Passive environment

The **overall system flow of the Financial Statement Analytics Pipeline** begins with collecting financial information from multiple sources such as **PDF reports, annual reports, XBRL filings, Excel/CSV files, web sources, APIs, and databases**. The collected data is passed to the **Data Ingestion and Extraction Layer**, where documents are processed and relevant text, tables, and financial fields are extracted using tools such as PyMuPDF, pdfplumber, Camelot, Tesseract OCR, and Arelle. The extracted information then moves to the **Data Processing and Standardization Layer**, where data cleaning, validation, standardization, entity mapping, and quality checks are performed. The processed information is stored in **PostgreSQL/MySQL** along with

financial fact tables, company information, statement metadata, audit records, and data lineage. The stored data is then passed to the **Analytics and Intelligence Engine**, which performs financial calculations, ratio and trend analysis, anomaly and risk detection, financial distress forecasting, What-If scenario simulation, and Explainable AI analysis using techniques such as Scikit-learn, XGBoost, SHAP, and LIME. The generated results are delivered through the **User Interface and Visualization Layer** using dashboards, reports, insights, alerts, notifications, and the What-If Simulator. The system supports both an **active environment**, where users upload documents and interact with the analysis in real time, and a **passive environment**, where scheduled processes automatically collect and update financial data. Finally, investors, analysts, auditors, and management can access the generated financial insights for informed decision-making, while infrastructure components such as **Docker**, **cloud/on-premise deployment**, **access control**, **encryption**, **monitoring**, **logging**, and **backup and recovery** support the reliability and security of the complete pipeline.



Dia 3.3.1 (c) Overall system flow

CHAPTER 4

PROJECT DESCRIPTION

4.1 METHODOLOGY

The **Financial Statement Analytics Pipeline** follows a systematic methodology that transforms raw financial documents into structured financial information and actionable insights. The methodology begins with uploading or collecting financial statements from sources such as **PDF reports, XBRL filings, Excel/CSV files, annual reports, APIs, and databases**. The required financial information is extracted using document-processing and OCR techniques and is then subjected to cleaning, validation, and standardization. Standardized financial records are stored in a structured database so that they can be reused for analysis without repeating the extraction process. The stored data is then processed for financial ratio calculation, trend analysis, financial health assessment, anomaly and risk detection, and forecasting. The system additionally provides a **What-If Scenario Simulator**, where users can modify variables such as revenue, expenses, and debt to study their possible effect on financial performance. Explainable AI techniques are used to identify the factors contributing to predictions, risks, anomalies, and changes in performance. Finally, the results are presented through an interactive dashboard containing KPIs, ratios, trends, risk indicators, scenario comparisons, and AI-generated explanations. This overall methodology integrates data engineering, financial analytics, machine learning, and visualization into a single workflow.

4.2 MODULES

The proposed system is divided into a set of interconnected modules. Each module performs a specific operation in the financial analytics workflow, while the output of one module becomes the input for the subsequent stage. The major modules are **Data Ingestion and Extraction, Data Cleaning and Standardization, Data Storage, Financial Analytics, Risk and Anomaly Detection, What-If Scenario Simulation, Explainable AI and Insights, and Interactive Dashboard**. This modular structure allows additional financial metrics, statement types, analytical models, and visualization components to be incorporated without redesigning the complete pipeline.

4.2.1 DATA INGESTION AND EXTRACTION MODULE

The Data Ingestion and Extraction Module forms the entry point of the system. It accepts financial information from **PDF documents, XBRL filings, Excel/CSV files, annual reports, APIs, and other supported sources**. PDF documents are processed to identify relevant text, tables, and financial values, while OCR is used when information is available in scanned documents. XBRL data is parsed to identify standardized financial facts and their associated tags. The extracted information is converted into an internal structured format containing details such as company, reporting period, statement type, financial metric, and value. This module reduces the need for repetitive manual data entry and prepares the information for subsequent processing.

4.2.2 DATA CLEANING AND STANDARDIZATION MODULE

The Data Cleaning and Standardization Module ensures that the extracted financial information is suitable for analysis. It identifies and handles **missing values, duplicate records, invalid values, inconsistent numerical formats, and variations in financial terminology**. Financial metrics having different names across documents are mapped to common standardized names, while reporting periods are converted into a consistent format for chronological comparison. Validation is performed before analytical calculations so that incorrect or inconsistent information does not affect the results. Pandas and NumPy can be used to perform these preprocessing and transformation operations.

4.2.3 DATA STORAGE MODULE

The Data Storage Module maintains the cleaned and standardized financial information in a structured database. **PostgreSQL or MySQL** can be used for persistent storage, with SQLAlchemy providing database connectivity. The module stores financial records, company information, statement metadata, reporting periods, and analytical data. Separating the stored information from the original input documents allows the system to retrieve previously processed records without performing extraction again. This also provides a foundation for historical analysis and comparison across multiple reporting periods.

4.2.4 FINANCIAL RATIO AND TREND ANALYSIS MODULE

The Financial Ratio and Trend Analysis Module uses standardized financial data to calculate important financial indicators. The module evaluates areas such as **liquidity, profitability, solvency, and efficiency** through appropriate

financial ratios. It also compares financial performance across different reporting periods to identify increases, decreases, and significant trends. The results help users understand the financial position and performance of an organization without manually performing repeated calculations.

4.2.5 FINANCIAL HEALTH AND RISK ANALYSIS MODULE

The Financial Health and Risk Analysis Module evaluates the overall condition of an organization using financial indicators and analytical models. The proposed system generates a **Financial Health Score from 0–100** based on relevant financial indicators and identifies unusual financial patterns. Machine-learning techniques such as **Scikit-learn and XGBoost** can be used for anomaly detection, risk assessment, and financial distress prediction. The module can also identify inconsistencies between different financial statements, such as changes in reported profit that are not supported by corresponding cash-flow movements.

4.2.6 WHAT-IF SCENARIO SIMULATION MODULE

The What-If Scenario Simulation Module enables users to examine hypothetical financial situations without changing the original stored financial data. Users can modify important parameters such as **revenue, expenses, and debt**, after which the system recalculates the relevant financial metrics. The resulting scenario is compared with the original financial position to show the possible effect of the changes. This module supports decision-making by allowing users to examine potential financial outcomes before implementing a particular strategy.

4.2.7 EXPLAINABLE AI AND INSIGHTS MODULE

The Explainable AI and Insights Module improves the interpretability of machine-learning-based financial analysis. **SHAP and LIME** can be used to identify the financial factors that contribute to a particular prediction, risk score, or anomaly. Instead of providing only a numerical prediction, the module generates understandable explanations regarding the factors responsible for the observed result. The system can then convert the analytical results into meaningful financial insights that assist users in interpreting risks, trends, and performance changes.

4.2.8 INTERACTIVE DASHBOARD MODULE

The Interactive Dashboard Module acts as the primary output interface of the system. It presents **KPIs, financial ratios, historical trends, financial health scores, risk indicators, anomalies, forecasts, scenario comparisons, and AI-generated explanations** through interactive visualizations. Technologies such as **React/Streamlit and Plotly** can be used to implement the dashboard, while Power BI can be used for business-intelligence reporting. The dashboard allows investors, analysts, auditors, and management to access the processed information in a clear and decision-oriented format.

4.2.9 REPORT GENERATION MODULE

The Report Generation Module consolidates the results produced by the different analytical components into a comprehensive financial analysis report. It can include the company's key financial indicators, calculated ratios, trends, identified risks and anomalies, What-If scenario results, and explanations generated by the AI component. The report can be made available for viewing or downloading, providing users with a permanent summary of the analysis.

4.2.10 PIPELINE ORCHESTRATION MODULE

The Pipeline Orchestration Module coordinates the different stages of the analytics workflow. **Apache Airflow** can be used to schedule recurring data extraction, preprocessing, analysis, and update operations. This is particularly useful when financial information needs to be processed periodically. The modular architecture also makes it possible to add new statement types, ratios, analytical models, and visualization components without redesigning the entire pipeline.

4.2.11. Evaluation Metrics

Data Extraction Accuracy

$$\text{Extraction Accuracy} = \frac{\text{Correctly Extracted Fields}}{\text{Total Fields}} \times 100$$

Precision

$$\text{Precision} = \frac{TP}{TP + FP}$$

Recall

$$\text{Recall} = \frac{TP}{TP + FN}$$

F1-Score

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Data Validation Accuracy

$$\text{Validation Accuracy} = \frac{\text{Correctly Validated Records}}{\text{Total Records}} \times 100$$

Standardization Accuracy

$$\text{Standardization Accuracy} = \frac{\text{Correctly Standardized Fields}}{\text{Total Fields}} \times 100$$

Data Completeness

$$\text{Data Completeness} = \frac{\text{Available Required Fields}}{\text{Total Required Fields}} \times 100$$

Data Consistency

$$\text{Data Consistency} = \frac{\text{Consistent Records}}{\text{Total Records}} \times 100$$

Ratio Calculation Accuracy

$$\text{Ratio Accuracy} = \frac{\text{Correctly Calculated Ratios}}{\text{Total Ratios}} \times 100$$

Percentage Error

$$\text{Percentage Error} = \frac{|\text{Actual Value} - \text{Calculated Value}|}{|\text{Actual Value}|} \times 100$$

Mean Absolute Error (MAE)

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

Root Mean Square Error (RMSE)

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

R² Score

$$R^2 = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$$

Anomaly Detection Accuracy

$$\text{Anomaly Accuracy} = \frac{\text{Correctly Detected Anomalies}}{\text{Total Anomalies}} \times 100$$

Risk Detection Precision

$$\text{Risk Precision} = \frac{TP}{TP + FP}$$

Risk Detection Recall

$$\text{Risk Recall} = \frac{TP}{TP + FN}$$

Risk Detection F1-Score

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Scenario Simulation Accuracy

$$\text{Scenario Accuracy} = \frac{\text{Correct Scenario Results}}{\text{Total Scenario Results}} \times 100$$

Scenario Error

$$\text{Scenario Error} = \frac{|\text{Reference Value} - \text{Scenario Value}|}{|\text{Reference Value}|} \times 100$$

System Response Time

$$\text{Response Time} = \text{End Time} - \text{Start Time}$$

Processing Throughput

$$\text{Throughput} = \frac{\text{Number of Records Processed}}{\text{Processing Time}}$$

Pipeline Reliability

$$\text{Reliability} = \frac{\text{Successful Executions}}{\text{Total Executions}} \times 100$$

Overall Project Description


Overall, the project combines **financial document processing, ETL, data validation, database management, financial calculations, machine learning, What-If simulation, Explainable AI, and data visualization** into one integrated pipeline. The workflow begins with financial-data ingestion and ends with decision-ready financial insights, reducing repetitive manual operations while improving consistency and accessibility of financial information. The project paper specifically describes the pipeline as connecting **financial-data extraction, transformation, storage, and analysis** into a unified process.

CHAPTER 5

IMPLEMENTATION AND RESULT DISCUSSION


5.1 Implementation

The Financial Statement Analytics Pipeline is designed as a modular web-based application that automates the processing and analysis of financial statement data. The system follows a sequential pipeline in which financial data is uploaded, extracted, cleaned, standardized, stored, and then used for financial analysis. The architecture separates the frontend interface, backend processing logic, and database storage so that each stage can be developed and extended independently.




The **Sign In module** provides secure user authentication for the LedgerSignal Financial Statement Analytics Pipeline. It allows registered users to access their saved financial analyses, financial health scores, and generated reports using their email address and password. The interface provides separate options for existing users to sign in and new users to create an account. This authentication mechanism ensures that user-specific financial analysis data is accessed only by authorized users and provides a personalized entry point to the analytics platform.

Home Page Module



The Home Page acts as the main introduction and navigation point of the LedgerSignal system. It presents the purpose of the application as a pipeline that transforms raw financial statements into meaningful financial insights. The page provides options such as starting an analysis, uploading financial data, opening the dashboard, and viewing the processing pipeline. The workflow displayed on the page follows the sequence of upload, extraction, validation, cleaning, standardization, ratio calculation, analysis, and insight generation. This provides users with a clear understanding of how the uploaded financial data is processed throughout the system.


Data Upload Module



The Data Upload module is used to provide financial statement data to the analytics pipeline. The interface supports uploading financial statement files in

CSV format and text-based PDF documents, with an XBRL-ready schema also indicated for future data compatibility. After a file is uploaded, the system automatically performs extraction, validation, cleaning, standardization, storage, ratio calculation, analysis, and insight generation. The implementation displays the progress of each processing stage, allowing the user to verify whether the uploaded data has successfully passed through the pipeline. The result of this module is a processed dataset that can be used by all subsequent analytical modules.

Data Quality Module




Stored canonical dataset (15 values)
Original labels are preserved alongside standardized names (traceability, spec 329).

Period	Statement	Standard metric	Original label	Value	Page	Method
2022	Balance Sheet	Equity	---	69,833	---	CSV-WD01
2022	Balance Sheet	Total Assets	Total_Assets	4,16,963	---	CSV-WD01
2022	Balance Sheet	Total Liabilities	Total_Liabilities	3,56,230	---	CSV-WD01
2022	Cash Flow Statement	Operating Cash Flow	Operating_Cash_Flow	42,600	---	CSV-WD01
2022	Income Statement	Net Income	Net_Income	2,217	---	CSV-WD01
2023	Balance Sheet	Equity	---	83,000	---	CSV-WD02
2023	Balance Sheet	Total Assets	Total_Assets	4,65,000	---	CSV-WD02
2023	Balance Sheet	Total Liabilities	Total_Liabilities	3,82,000	---	CSV-WD02
2023	Cash Flow Statement	Operating Cash Flow	Operating_Cash_Flow	52,000	---	CSV-WD02


The Data Quality module evaluates the quality and reliability of the extracted financial data before it is used for analysis. The implementation performs validation checks such as duplicate detection, value validation, consistency checking, sign checking, balance-sheet validation, outlier identification, and unit-related checks. Conflicting duplicate records are displayed as warnings instead of being silently modified, allowing the user to review the original source. The module also displays completeness, validity, consistency, extraction coverage, and standardization information. The stored canonical dataset preserves the standardized financial metrics along with their original labels, which improves traceability and supports reliable downstream analysis.


Financial Statements Module



The Financial Statements module presents the processed financial information in a structured form. The extracted data is organized into major financial statement categories such as the Income Statement, Balance Sheet, and Cash Flow Statement. Financial metrics are displayed across different reporting periods, allowing users to compare values between years. The implementation also highlights changes between periods so that variations in important financial indicators can be easily observed. This module provides the standardized financial data required for ratio calculation, trend analysis, financial health evaluation, and risk assessment.

Ratios Module







The Ratios module automatically calculates financial ratios from the standardized financial statement data. The implementation groups ratios into categories such as liquidity, profitability, solvency, efficiency, and growth. Each ratio displays its formula, required input values, calculated result, previous-period comparison, and availability status. When the required financial information is unavailable, the system displays a clear indication that there is not enough data instead of estimating or substituting a value. This ensures that


the financial ratio analysis remains transparent and based only on the available processed data.


Trends and Patterns Module



The Trends and Patterns module analyzes financial metrics across multiple reporting periods to identify changes and recurring patterns. The implementation provides graphical representations for indicators such as Net Income, Operating Cash Flow, Total Assets, Total Liabilities, and Equity. Users can select different metrics and comparison periods to examine financial performance over time. The system also identifies patterns such as continuous increases in financial metrics and presents them as positive or relevant observations. This module helps users understand the direction and development of the company's financial performance rather than examining individual values independently.


Financial Health Score Module







The Financial Health Score module combines multiple financial indicators into a single interpretable measure of financial condition. The implementation considers components such as liquidity, profitability, solvency, growth, and cash flow and maps their values into component scores using documented scoring rules. The resulting health score is displayed together with a component-wise breakdown so that users can understand how each financial area contributes to the overall result. The system also provides score history across reporting periods, enabling comparison of financial health over time. When a required component cannot be calculated because of missing data, the system identifies the limitation rather than generating an estimated value.

Risk and Anomalies Module



The Risk and Anomalies module identifies unusual financial behaviour and potential risk indicators from the processed financial data. The implementation uses statistical anomaly detection and structural financial rules to examine changes in important metrics. Detected anomalies are displayed with their observed values, reference or typical values, detection method, severity, and supporting information. The module also identifies financial risk drivers such as high leverage, losses, or cash-flow-related concerns. The results are presented as indicators requiring further investigation rather than as direct accusations or conclusions, making the analysis more transparent and appropriate for financial decision-support.


Profit–Cash Flow Consistency Module



The Profit–Cash Flow Consistency module examines whether changes in reported profit are supported by corresponding changes in operating cash flow. The implementation compares Net Income and Operating Cash Flow across


available periods and calculates their relative movement. A consistency indicator and graphical comparison are provided to help users identify whether earnings are generally supported by cash generation. The underlying values and period-level findings are also displayed so that the user can trace the consistency result back to the financial data used by the system. This module provides an additional analytical view of the relationship between accounting profit and actual operating cash generation.

Financial Distress Forecast Module



The Financial Distress Forecast module provides a risk-oriented indication of potential financial distress based on the processed financial metrics. The current implementation uses a clearly labelled rule-based heuristic rather than presenting the result as a validated machine-learning prediction. The module considers financial indicators such as liability-to-asset ratio, debt-to-equity, return on assets, operating cash flow relative to liabilities, and other available metrics. Each feature is shown along with its contribution and direction toward the risk score. This implementation improves transparency by allowing users to understand why a particular risk level has been produced while clearly indicating the limitations of the heuristic approach.

What-If Simulator Module



The What-If Simulator enables users to examine how changes in financial assumptions may affect the company's financial condition. The implementation allows users to modify assumptions related to revenue, operating expenses, debt, interest expense, operating cash flow, and capital expenditure. After the assumptions are applied, the system recalculates the relevant financial values, ratios, and health score and compares the base case with the simulated scenario. The results are presented through comparative charts and risk indicators. The module clearly identifies the output as a scenario simulation rather than a prediction of actual future results, allowing users to study the potential impact of different financial conditions.

AI Copilot Module

The AI Copilot interface is a web-based application for financial analysis. It features a sidebar with navigation options: Overview, Data Upload, Data Quality, Financial Statements, Ratios, Trends, Health Score, Risk & Anomalies, Cash Flow Consistency, Forecast, What If Simulator, AI Copilot, and Reports / Exports. The main content area is titled 'AI Copilot' and includes a search bar, a list of suggested questions, and a 'Generated insights (from the pipeline)' section.

Generated insights (from the pipeline)
Produced automatically after processing — each cites its evidence.

Revenue increased by 12.4%
Revenue moved from 63799 (2024) to 71604 (2025), a change of 12.4%. This is calculated directly from the stored metrics below:
Revenue 2024: 638.4K Revenue 2025: 716.1K

Net income grew 31.1%
Net income moved from 59248 (2024) to 77670 (2025), a change of 31.1%.
Net income 2024: 59.2K Net income 2025: 77.7K

Net margin expanded 1.5 pp
Net margin was 9.3% in 2024 and 10.8% in 2025 (Net Income ÷ Revenue for each period).
Revenue 2024: 716.1K Net income 2024: 66.9K Revenue 2025: 716.1K Net income 2025: 77.7K

Potential anomaly: Net Income (2023)
Year-over-year change in Net Income (+1217.7% in 2023) falls outside the typical range across periods (Q2 rule). Unusual financial pattern — verify against the source. Underlying values: Net Income 2022 → 2023 → 2024 → 2025: 30420, Typical Net range (Q2) → 372.5% to 944.3%.
Net income 2023: 30.4K

Debt-to-equity stands at 0.99x
Total liabilities (405977) ÷ equity (411003) = 0.99x in 2025. This is within the reference band used by the health model's solvency component.

Suggested questions
Tap to ask — the copilot answers from your data only.

What changed this year?
Why is cash flow inconsistent with profit?
What are the major financial risks?
Which ratios deteriorated?
What happened to revenue?
Explain the anomaly
Why did the health score change?
Show the strongest positive trends

Generated insights (from the pipeline)
Produced automatically after processing — each cites its evidence.

Revenue increased by 12.4%
Revenue moved from 63799 (2024) to 71604 (2025), a change of 12.4%. This is calculated directly from the stored metrics below:
Revenue 2024: 638.4K Revenue 2025: 716.1K

Net income grew 31.1%
Net income moved from 59248 (2024) to 77670 (2025), a change of 31.1%.
Net income 2024: 59.2K Net income 2025: 77.7K

28 ratios deteriorated, 40 improved
Deteriorated: Current Ratio: 1.14 → 0.94; Quick Ratio: 0.91 → 0.72; Gross Margin: 42.03 → 43.81%; Net Profit Margin: 7.10 → -0.53%; Return on Assets: 7.93 → -0.59%; Return on Equity: 24.13 → -1.86%; Debt-to-Equity: 2.04 → 2.17; Debt-to-Assets: 0.67 → 0.68; Equity Ratio: 0.33 → 0.32; Interest Coverage: 13.75 → 5.17; Revenue Growth: 9.40 → 11.83%; Net Income Growth: -105.16 → 1217.74%; Asset Growth: 10.02 → 14.09%; Operating Cash Flow Growth: 0.92 → 81.69%; Gross Margin: 43.81 → 46.98%; Operating Margin: 2.38 → 6.41%; Receivables Turnover: 11.00 → 11.50; Asset Growth: 14.09 → 16.38%; Cash Ratio: 0.44 → 0.44; Gross Margin: 46.98 → 48.85%; Operating Margin: 6.41 → 10.75%; Revenue Growth: 10.99 → 12.38%; Asset Growth: 18.38 → 30.91%; Current Ratio: 1.06 → 1.05; Cash Ratio: 0.44 → 0.40; Gross Margin: 48.85 → 50.29%; Operating Margin: 10.75 → 11.16%; Return on Equity: 20.72 → 18.09%; Improved: Asset Turnover: 1.12 → 1.11; Receivables Turnover: 14.28 → 12.13; Cash Ratio: 0.25 → 0.35; Operating Margin: 5.30 → 2.38%; Asset Turnover: 1.11 → 1.09; Receivables Turnover: 12.13 → 11.00; Current Ratio: 0.94 → 1.05; Quick Ratio: 0.72 → 0.84; Cash Ratio: 0.35 → 0.44; Net Profit Margin: -0.53 → 5.29%; Return on Assets: -0.59 → -5.76%; Return on Equity: -1.86 → 15.07%; Debt-to-Equity: 2.17 → 1.61; Debt-to-Assets: 0.68 → 0.62; Equity Ratio: 0.32 → 0.38; Interest Coverage: 5.17 → 11.58; Asset Turnover: 1.09 → 1.02; Revenue Growth: 11.83 → 10.99%; Net Income Growth: 1217.74 → 94.73%; Operating Cash Flow Growth: 81.69 → 36.41%; Current Ratio: 1.05 → 1.06; Quick Ratio: 0.84 → 0.87; Net Profit Margin: 5.29 → 9.29%; Return on Assets: 5.76 → 9.48%; Return on Equity: 15.07 → 20.72%; Debt-to-Equity: 1.61 → 1.19; Debt-to-Assets: 0.62 → 0.54; Equity Ratio: 0.38 → 0.46; Interest Coverage: 11.58 → 28.51; Asset Turnover: 1.02 → 0.88; Receivables Turnover: 11.50 → 13.59; Net Income Growth: 94.73 → 31.02%; Operating Cash Flow Growth: 36.41 → 20.40%; Quick Ratio: 0.87 → 0.82; Net Profit Margin: 9.29 → 10.81%; Return on Assets: 9.48 → 9.49%; Debt-to-Equity: 1.19 → 0.99; Debt-to-Assets: 0.54 → 0.50; Equity Ratio: 0.46 → 0.50; Interest Coverage: 28.51 → 35.17.

Question: "Which ratios deteriorated?" Answered from analysis DBs. Rule-based engine over stored data (the external LLM).

Generated insights (from the pipeline)
Produced automatically after processing — each cites its evidence.

Revenue increased by 12.4%
Revenue moved from 63799 (2024) to 71604 (2025), a change of 12.4%. This is calculated directly from the stored metrics below:
Revenue 2024: 638.4K Revenue 2025: 716.1K

Net income grew 31.1%
Net income moved from 59248 (2024) to 77670 (2025), a change of 31.1%.
Net income 2024: 59.2K Net income 2025: 77.7K


Suggested questions
Tap to ask — the copilot answers from your data only.

What changed this year?
Why is cash flow inconsistent with profit?
What are the major financial risks?
Which ratios deteriorated?
What happened to revenue?
Explain the anomaly
Why did the health score change?
Show the strongest positive trends

The AI Copilot provides an interactive interface for asking questions about the processed financial analysis. The implementation generates responses using the stored analysis results and provides evidence associated with the metrics used to produce each response. Users can ask questions related to anomalies, financial changes, cash-flow consistency, financial risks, deteriorating ratios, revenue changes, health-score variations, and positive trends. Generated insights are also displayed automatically after processing. This module improves interpretability

by converting analytical results into understandable explanations while maintaining a connection between the generated insight and the underlying financial metrics.

Reports and Exports Module



The Reports and Exports module provides downloadable outputs from the completed financial analysis. The implementation generates a clean financial dataset, ratio report, risk report, and complete analysis summary. The exported results are generated from the stored analysis data so that the information presented in the reports remains consistent with the values used in the dashboard. The complete summary includes financial data quality, ratios, health

analysis, risk findings, cash-flow consistency, forecast information, insights, and limitations. This module enables users to preserve and share the results of the financial analysis in a structured format.

Dashboard and Overall Analysis Module

The Dashboard integrates the outputs generated by the individual modules into a consolidated financial analysis view. It presents important financial indicators, trends, financial health information, risk and anomaly findings, cash-flow consistency, and key insights in a single interface. The dashboard allows users to move between detailed modules whenever additional investigation is required. The implementation demonstrates the complete flow of the proposed financial statement analytics pipeline, from raw data ingestion to standardized financial information, analytical calculations, risk identification, scenario analysis, explainable insights, and downloadable reports. Overall, the implemented system provides an integrated environment for transforming financial statement data into structured and interpretable analytical results.

5.2.DISCUSSION

The implementation of the Financial Statement Analytics Pipeline demonstrates the effectiveness of integrating financial data processing, validation, analysis, risk assessment, and reporting into a single system. The developed LedgerSignal application provides a systematic flow in which financial statement data is uploaded, extracted, validated, cleaned, standardized, stored, and subsequently used for financial calculations and analytical operations. This reduces the need for repeated manual processing and provides a consistent structure for analysing financial information.

The Data Quality module plays an important role in improving the reliability of the analysis by identifying duplicate records, conflicting values, missing information, and other validation issues before the data is used for further calculations. The standardized financial statements allow important metrics such as Net Income, Total Assets, Total Liabilities, Equity, and Operating Cash Flow to be compared across reporting periods. The Ratios and Trends modules further transform these values into meaningful financial indicators and visual

patterns, making it easier to understand changes in the company's financial performance.

The Financial Health Score provides a consolidated representation of the company's financial condition by considering different financial components such as profitability, solvency, growth, and cash flow. The Risk and Anomalies module complements this analysis by identifying unusual changes and potential financial risk factors. The Profit–Cash Flow Consistency module provides an additional perspective by comparing reported profit with operating cash generation. These combined analytical components provide a broader understanding of financial performance rather than depending on a single financial indicator.

The What-If Simulator demonstrates the usefulness of scenario-based analysis by allowing users to modify important financial assumptions and observe their effect on financial metrics and the overall health score. This helps users understand how changes in revenue, expenses, debt, interest, cash flow, and capital expenditure can influence the financial condition of the organization. The Financial Distress Forecast and AI Copilot further improve interpretation by presenting risk indicators and evidence-based explanations derived from the processed financial data.

The Reports and Exports module provides a practical outcome of the complete pipeline by allowing processed datasets, ratio results, risk information, and analysis summaries to be exported. Overall, the implementation shows that the proposed system can convert raw financial statement data into structured, traceable, and interpretable financial information. The results indicate that integrating data quality checking, financial analysis, risk identification, scenario simulation, and explainable insights in one pipeline can provide a more efficient and transparent approach to financial statement analysis.

CHAPTER 6

CONCLUSION AND FUTURE WORK

6.1 Conclusion

The **Financial Statement Analytics Pipeline**, implemented as **LedgerSignal**, provides an integrated web-based platform for transforming raw financial statement data into structured, analysed, and interpretable financial information. The system was designed to address the limitations of conventional financial analysis, where financial data is often extracted manually, cleaned using spreadsheets, analysed through repeated calculations, and interpreted separately across different tools. By combining these activities into a single pipeline, the developed system provides a systematic approach to financial statement processing and analysis.

The implemented system begins with the **Data Upload module**, where financial statement files can be provided in supported formats. Once the data is uploaded, the pipeline automatically proceeds through extraction, validation, cleaning, standardization, storage, and analysis. This provides a continuous workflow instead of requiring the user to manually transfer financial values between different applications. The processing status is displayed to the user, making the individual stages of the pipeline visible and easier to understand.

The **Data Quality module** plays an important role in ensuring that the extracted information is suitable for further analysis. The system performs validation

checks for duplicate records, conflicting values, missing information, unusual values, and other data-quality conditions. Instead of silently modifying questionable data, the system displays warnings and preserves traceability to the extracted information. The standardized canonical dataset allows financial metrics to be represented using common names and structures, which makes them easier to use consistently throughout the subsequent modules.

The **Financial Statements module** organizes the processed information into major financial statement categories, including the Income Statement, Balance Sheet, and Cash Flow Statement. The financial values are presented across different reporting periods, allowing users to observe changes in important financial metrics. This structured representation acts as the foundation for the analytical modules and reduces the complexity of working directly with raw financial statement data.

The **Ratios module** provides automated calculation of financial ratios based on the available standardized metrics. The system groups the ratios into appropriate financial categories and displays the formula and required inputs along with the calculated result. An important feature of the implementation is that the system does not substitute unavailable values with estimated values. When the required inputs are missing, the corresponding ratio is marked as unavailable. This improves transparency and prevents unsupported calculations from being presented as actual financial results.

The **Trends and Patterns module** extends the analysis from individual financial values to changes across reporting periods. Graphical representations of Net Income, Operating Cash Flow, Total Assets, Total Liabilities, and Equity allow users to identify increasing or decreasing patterns. The module also provides automatically detected patterns based on the available financial series. This makes it easier to understand the direction of financial performance and identify significant changes that may require further examination.

The **Financial Health Score module** provides a consolidated representation of the financial condition of the organization. The score is derived from different financial components such as profitability, solvency, growth, and cash flow. The component breakdown allows users to understand the contribution of individual financial areas rather than receiving only a single overall score. The implementation also maintains a score history so that financial health can be compared across different periods. Where a component cannot be calculated

because the required financial information is unavailable, the system identifies the limitation and adjusts the scoring process accordingly.

The **Risk and Anomalies module** provides an additional level of financial analysis by identifying unusual financial behaviour and potential risk indicators. The system examines financial values and changes across periods and presents detected anomalies together with their observed values, reference values, detection method, severity, and supporting information. This provides users with a structured method for identifying financial conditions that may require further investigation. The module is designed as an analytical indicator rather than a mechanism for making unsupported claims about financial misconduct.

The **Profit–Cash Flow Consistency module** provides a separate assessment of the relationship between reported profit and operating cash flow. By comparing Net Income with Operating Cash Flow across different periods, the system determines whether changes in reported earnings are generally supported by operating cash generation. The module provides both a consistency indicator and the underlying period-level values, making the result traceable to the financial data used in the calculation.

The **Financial Distress Forecast module** provides a risk-oriented assessment using the financial indicators available in the system. The current implementation explicitly identifies the distress assessment as a **rule-based heuristic** rather than a validated machine-learning prediction. This distinction is important because a reliable predictive model would require a sufficiently large labelled dataset containing historical examples of financial distress. The module nevertheless provides transparency by displaying the financial features used, their contribution to the risk score, and the direction in which they influence the assessment.

The **What-If Simulator** adds scenario-based analysis to the pipeline. Users can modify assumptions related to revenue, operating expenses, debt, interest expense, operating cash flow, and capital expenditure and then observe the resulting changes in financial metrics and financial health. The system compares the base condition with the simulated scenario through numerical and graphical representations. This allows users to understand the possible effect of changes in financial assumptions without altering the original stored financial data. The system also clearly distinguishes scenario simulation from actual forecasting.

The **AI Copilot module** improves the accessibility and interpretation of the financial analysis. Users can ask questions about anomalies, financial changes, cash-flow consistency, financial risks, deteriorating ratios, revenue changes, health-score variations, and positive trends. The responses are derived from the processed analytical artifacts and are associated with the stored financial metrics used to generate the explanation. This provides a more understandable interface for users who may not want to manually inspect every analytical table or chart.

The **Reports and Exports module** completes the analytical workflow by allowing users to export the processed financial information and analytical results. The implementation provides outputs such as the clean financial dataset, ratio report, risk report, and complete analysis summary. Since the reports are generated from the stored analysis results, the exported information is intended to remain consistent with the results displayed within the application.

The implemented system also demonstrates the importance of **data transparency and traceability** in financial analytics. Financial values used by the different modules can be traced back to the processed dataset, while warnings and unavailable values are explicitly communicated to the user. This approach is particularly useful in financial applications because unsupported assumptions or silently modified values could lead to incorrect interpretations.

However, the implementation also identifies areas where further development is required. Some financial metrics may not be available in the input dataset, which can prevent certain ratios or health-score components from being calculated. The current financial distress assessment is a heuristic rather than a trained predictive model, and the What-If results represent simulations rather than actual future forecasts. These limitations are explicitly communicated within the application and provide a clear direction for future development.

Overall, the project demonstrates that an integrated financial analytics pipeline can significantly simplify the process of converting raw financial statements into structured analytical information. By combining data quality management, financial statement organization, ratio analysis, trend analysis, financial health scoring, anomaly detection, cash-flow consistency analysis, scenario simulation, explainable insights, and reporting within one platform, LedgerSignal provides a comprehensive foundation for financial analysis and decision support.

6.2 Future Work

Although the current implementation provides the major components of the proposed Financial Statement Analytics Pipeline, several improvements can be introduced to increase its accuracy, scalability, automation, and practical applicability.

Advanced Financial Document Extraction

Future versions can improve the document extraction process by supporting a wider range of financial documents and complex document structures. Financial statements may contain tables with multiple columns, merged cells, footnotes, images, different units, and non-standard layouts. Advanced document-processing techniques can be incorporated to improve the extraction of financial metrics from such documents.

OCR technology can also be integrated to process scanned financial statements and image-based documents. This would extend the system beyond text-based PDFs and CSV files and make the pipeline more suitable for real-world financial documents that are not available in machine-readable form.

XBRL-Based Data Integration

The existing system indicates an XBRL-ready schema, which can be further developed into complete XBRL processing. Future work can include automatic extraction of standardized financial concepts from XBRL filings, mapping XBRL tags to the system's canonical financial metrics, and handling different reporting periods and units automatically. This can reduce manual mapping requirements and improve the consistency of financial data obtained from structured filings.

Improved Data Cleaning and Validation

The data-quality component can be extended with more advanced validation mechanisms. Future versions can identify inconsistencies between related financial statements, validate accounting relationships, detect incorrect units, and automatically identify potential data-entry errors. For example, relationships between assets, liabilities, and equity can be checked more extensively.

Duplicate handling can also be improved by introducing source-priority rules, confidence levels, and user-assisted verification when multiple extracted values

conflict. This would help the system handle large and complex financial datasets more reliably.

Advanced Financial Ratio Analysis

The current ratio module can be expanded to include a larger set of financial indicators covering liquidity, profitability, solvency, efficiency, valuation, and growth. Additional ratios can provide a more comprehensive assessment of an organization's financial condition.

Future versions can also provide industry-specific ratio interpretation. Instead of evaluating ratios independently, the system could compare them against historical company performance, industry benchmarks, or selected competitors. This would provide additional context for determining whether a ratio represents a strong or weak financial position.

Machine Learning-Based Risk Prediction

One of the major areas for future development is the replacement or enhancement of the current rule-based financial distress indicator with a properly trained machine learning model. Models such as Logistic Regression, Random Forest, XGBoost, or other suitable classification techniques can be evaluated using historical financial datasets containing labelled examples of financially healthy and distressed organizations.

A sufficiently large and representative dataset would be required for this development. The model could use features such as profitability ratios, leverage ratios, liquidity indicators, cash-flow measures, growth rates, and other financial characteristics. The performance of the model could then be evaluated using appropriate classification metrics such as precision, recall, F1-score, ROC-AUC, and confusion-matrix analysis.

Explainable Machine Learning

Once a validated machine learning model is introduced, explainability techniques can be incorporated to show why a particular financial risk prediction was generated. Methods such as SHAP can be used to identify the financial features that contribute most strongly to an individual prediction.

This would improve the current explainability framework by moving from rule-based feature contributions toward model-based explanations. Users would be

able to understand not only the predicted risk level but also which financial indicators were responsible for increasing or decreasing that risk.

Improved Financial Forecasting

The existing forecasting component can be developed into a more comprehensive forecasting system by using larger historical datasets. Future versions can evaluate appropriate forecasting techniques based on the type and availability of financial time-series data.

Forecasting can be extended to variables such as revenue, operating expenses, net income, operating cash flow, liabilities, and other important financial metrics. Different forecasting models can be compared using evaluation metrics such as MAE, RMSE, and other suitable measures before selecting the most appropriate model.

Enhanced What-If Scenario Analysis

The What-If Simulator can be extended to support multiple simultaneous scenarios and more complex relationships between financial variables. Users could create different scenarios such as optimistic, normal, and adverse conditions and compare their effects side by side.

Future versions could also support scenario saving, scenario history, sensitivity analysis, and automatic identification of variables that have the greatest influence on financial health. This would make the simulator more useful for financial planning and strategic analysis.

Industry and Peer Benchmarking

A future version can introduce benchmarking functionality to compare the analyzed organization with industry averages or selected peer organizations. Such comparisons could be performed for profitability, solvency, liquidity, growth, cash flow, and other financial indicators.

Benchmarking would provide additional context because an individual ratio may not be meaningful without knowing how it compares with similar organizations. This feature could help users identify areas where the organization performs above or below industry expectations.

Enhanced AI Copilot

The AI Copilot can be expanded to support more advanced financial questions and multi-step analytical queries. Instead of answering only individual questions, the system could generate structured financial summaries covering performance, risks, trends, cash flow, and major changes.

Future development can also include comparison-based questions, such as comparing two reporting periods, identifying the largest changes, explaining why a health score changed, and summarizing the main financial strengths and weaknesses. The evidence-based approach used in the current implementation can be retained to ensure that responses remain connected to the stored financial data.

Automated Report Generation

The reporting module can be improved by introducing professionally formatted PDF reports in addition to the existing downloadable data and text-based outputs. Future reports can include charts, financial statements, ratio tables, health-score summaries, risk findings, scenario comparisons, and AI-generated explanations.

Reports could also be generated automatically after every completed analysis and stored with the corresponding analysis record. This would make it easier for users to maintain a historical record of financial analyses.

Cloud Deployment and Scalability

The current web application can be further developed into a scalable cloud-based platform. A cloud deployment would allow multiple users to access the system remotely and process larger financial datasets. Cloud-based databases and object storage can be used for managing uploaded documents and processed financial information.

The backend can also be containerized and deployed using scalable services so that computationally intensive tasks such as document extraction, data processing, forecasting, and machine learning can be handled efficiently.

Security and Multi-User Access

Future work can strengthen the authentication and authorization mechanisms to support multiple users and organizations. Role-based access control can be introduced so that administrators, analysts, and other users have appropriate levels of access.

Additional security mechanisms such as secure password handling, session management, encrypted data transmission, access logging, and controlled document storage can be incorporated to protect sensitive financial information.

Real-Time Financial Data Integration

The pipeline can eventually be connected to external financial data sources and APIs to reduce the dependency on manually uploaded files. This could allow financial information to be periodically updated and automatically processed by the analytics pipeline.

Automated data refresh mechanisms could also trigger recalculation of ratios, trends, health scores, risks, and insights whenever new financial data becomes available.

Improved User Interface and Visualization

The dashboard can be enhanced with more interactive visualizations, customizable financial KPI cards, advanced filtering, comparison controls, and drill-down functionality. Users could select specific financial periods, metrics, or statement categories and dynamically explore the underlying data.

Visualization improvements can also make complex analytical findings easier to understand for non-technical users while retaining detailed information for financial analysts.

Comprehensive Testing and Evaluation

Future work should include extensive testing using multiple companies, financial statement formats, reporting periods, and data-quality conditions. The current implementation demonstrates the pipeline using available financial data, but broader testing would help determine how well the system performs across different types of inputs.

Separate evaluation can be conducted for extraction accuracy, data validation, standardization, ratio calculation, anomaly detection, forecasting, scenario simulation, system response time, and overall pipeline reliability. This would provide stronger quantitative evidence for the performance of the complete system.

Overall Future Direction

The long-term objective is to evolve LedgerSignal from an academic financial analytics application into a more robust and scalable financial decision-support platform. The integration of advanced document intelligence, standardized financial data processing, machine learning-based risk prediction, explainable AI, forecasting, benchmarking, scenario analysis, cloud infrastructure, and automated reporting can significantly extend the capabilities of the current system.

Thus, the future development of the Financial Statement Analytics Pipeline can focus on **improving data extraction accuracy, strengthening analytical models, increasing automation, supporting larger datasets, improving explainability, and providing more reliable financial decision-support capabilities** while maintaining the transparency and traceability established in the current implementation.

REFERENCES

1. Financial Big Data Analysis and Early Warning Platform: A Case Study

Authors: Liang et al.

Year: 2020

Focus: Big-data platform for financial risk analysis and early warning using heterogeneous data collection, data cleaning, data fusion, cloud computing, and AI.

2. Intelligent Fraud Detection in Financial Statements Using Machine Learning and Data Mining

Authors: Ashtiani & Raahemi

Year: 2021

Focus: Review of machine-learning and data-mining approaches for detecting fraudulent financial statements. The study reviews 47 research papers and emphasizes anomaly detection and advanced ML techniques.

3. Financial Knowledge Graph Based Financial Report Query System

Authors: Zehra et al.

Year: 2021

Focus: Uses financial knowledge graphs and ontology-based information extraction to extract information from annual financial reports and support financial queries.

4. Sentiment Analysis in Finance: From Transformers Back to eXplainable Lexicons (XLex)

Authors: Rizinski et al.

Year: 2024

Focus: Combines transformer models with SHAP explainability for financial sentiment analysis, improving financial vocabulary coverage, classification accuracy, efficiency, and interpretability.

5. AI Driven Fraud Detection Models in Financial Networks: A Comprehensive Systematic Review

Authors: Sarna et al.

Year: 2025

Focus: Reviews AI, machine learning, deep learning, hybrid models, and graph

neural networks for financial fraud detection, while identifying explainability, privacy, bias, and scalability as challenges.

6. Multimodal Fraud Detection in Financial Statements: A Trimodal Attention Network With Contrastive Evidence Chain Construction

Authors: Ping et al.

Year: 2026

Focus: Proposes TriAudit, combining financial ratios, textual information, and document-layout features for financial-statement fraud detection with interpretable evidence.

7. Mining Corporate Annual Reports for Intelligent Detection of Financial Statement Fraud – A Comparative Study of Machine Learning Methods

Authors: Hajek & Henriques

Year: 2017

Focus: Combines financial information and managerial comments with machine-learning methods for financial-statement fraud detection.

8. A Research on Text Information Extraction from Annual Report Based on Domain Ontology

Year: 2017

Focus: Applies domain ontology and NLP techniques to extract information from annual reports. It supports automated financial-document extraction and data standardization.

9. Segment Information Extraction from Financial Annual Reports Using Neural Network

Authors: Ito, T., Sakaji, H., and Izumi, K.

Year: 2019

Focus: Uses neural-network techniques to extract segment information from financial annual reports and supports automated extraction of structured financial information.

10. FinBERT: Financial Sentiment Analysis with Pre-trained Language Models

Author: Araci, D.

Year: 2019

Focus: Introduces a BERT-based language model specifically designed for financial-domain sentiment analysis and financial-document understanding.

11. FinQA: A Dataset of Numerical Reasoning over Financial Data

Authors: Chen, Z. et al.

Year: 2021

Focus: Introduces a financial question-answering dataset requiring numerical reasoning over financial reports, supporting AI-based financial insight generation.

12. TAT-QA: A Question Answering Benchmark on a Hybrid of Tabular and Textual Content in Finance

Authors: Zhu, F. et al.

Year: 2021

Focus: Develops a benchmark combining tabular and textual financial information and requiring numerical reasoning, supporting the integration of financial tables and narrative information.

13. On the Determinants and Prediction of Corporate Financial Distress in India

Year: 2021

Focus: Investigates financial ratios and machine-learning models such as SVM and ANN for predicting corporate financial distress in the Indian context.

14. Predicting Financial Distress of Chinese Listed Companies Using Machine Learning: To What Extent Does Textual Disclosure Matter?

Year: 2023

Focus: Studies whether combining textual disclosures with financial information improves financial-distress prediction.

15. Assessing Large Language Models Used for Extracting Table Information from Annual Financial Reports

Year: 2024

Focus: Investigates the use of large language models for extracting and interpreting tabular information from annual financial reports. This directly supports modern AI/LLM-based financial-document and table extraction.

FINANCIAL STATEMENT ANALYTICS PIPELINE: AN AUTOMATED FRAMEWORK FOR FINANCIAL DATA EXTRACTION, STANDARDIZATION AND ANALYSIS

K M Padmasri
230701223@rajal
akshmi.edu.in
Rajalakshmi
Engineering
College ,Chennai

A K Pooja
230701226@rajal
akshmi.edu.in
Rajalakshmi
Engineering
College, Chennai

R Raghul
230701251@rajal
akshmi.edu.in
Rajalakshmi
Engineering
College, Chennai

Mrs. B .Vani
(Assistant
Professor)
Rajalakshmi
Engineering
College, Chennai

I . ABSTRACT

Financial statements are the primary source of information about an organization's financial position and performance, yet upon manually analyzing such reports, one usually spends a great deal of effort and time to draw any practical conclusions from the data. In this paper, the authors introduce a Financial Statement Analytics Pipeline that was designed as a student project to create a fully automated process for transforming raw financial data into useful analysis. The designed pipeline works with structured data from the financial statements, analyzes and validates the relevant attributes that were defined in the pipeline for the

current project, formats the data into a usable form, stores it in a database, derives various financial ratios, and conducts trend analysis reporting. The proposed system aims to minimize the amount of manual operations while ensuring greater consistency and availability of financial data. The project combines different aspects of data processing, database and financial analysis into a unified structure. The approach is developed for income statements, balance sheets and cash flow reports, and it allows defining common liquidity, profitability, solvency, and efficiency ratios.

II . INTRODUCTION

By means of financial statement analysis, a firm's commercial performance can be appreciated properly because this analysis is a crucial step for investors, executives, and creditors to get information about a company's profitability, liquidity, solvency, and efficiency. Studies show that financial statement analysis is a valuable tool for fundamental analysis and forecasting. In addition, it should be noted that ratio analysis is equally important since it helps to interpret the data provided. In addition, traditional methods of processing financial statements imply that data have to be collected and entered into a spreadsheet manually. However, various challenges might be faced while gathering financial data due to terminology differences along with missing values or variations in formatting. The contribution of software analytics procedure can solve this problem as the procedure assures that there is a clear and consistent system for collecting, cleaning, standardizing, and analyzing financial data. Consequently, this paper presents a financial analytics pipeline that links the data preparation phase with the analysis stage instead of the conventional approach to

financial ratio calculation that implies treating financial data as a separate stage. The literature survey reviews existing research on financial statement analysis, financial data extraction, fraud and anomaly detection, financial distress prediction, and AI-based financial reasoning. Previous studies have explored machine learning, NLP, financial ratios, and automated extraction techniques for analyzing financial information. Research such as FinQA and TATQA also highlights the importance of combining numerical and textual information for meaningful financial analysis. However, most existing studies focus on individual tasks rather than providing a complete end-to-end system. The proposed Financial Statement Analytics Pipeline addresses this gap by integrating data extraction, cleaning, standardization, financial analysis, risk detection, What-If simulation, Explainable AI, and interactive visualization into a unified platform.

III. PROBLEM STATEMENT

The primary issue dealt with in this project relates to the challenge of transforming raw financial statement data into

coherent and interpretable analytical information. Manual processes may include repetitive data entry, inconsistent naming conventions, mathematical mistakes, and the challenge of comparing several reporting periods. The purpose of the project is to construct a central hub, capable of automating these steps and preparing financial ratios and trends based on standardized information.

IV OBJECTIVES

This project intends to create an easy-to-use tool for entering and uploading financial reporting information along with deriving relevant financial metrics from the data provided. The emphasis will be on ensuring that the monetary information with respective IDs is accurate, consistent, and sanitized before storing the standardized data properly in the data base. The system is able to compute essential financial ratios and effectiveness indicators, as well as make a comparison of the profitability for different periods. Moreover, it represents the outcomes of financial analysis in a simple and understandable language, helping eliminate the unnecessary manual calculations

and improving efficiency of the financial analysis.

V. RELATED WORK


Financial statement analysis has been, for a long time, centered on information accumulation from accounting reports through ratio analysis, comparative analysis and forecasting. According to Trigueiros, ratios have a power to give useful information, albeit, at the same time, make it possible to lose information inherent in the ratios themselves prompting quest for more technology-driven solutions for analysis of accounting data. [3] Newer studies assess the idea of automated financial statement analysis. They confirm the capacity of automation to save humans' efforts and facilitate more efficient types of financial statement analysis. [4] The studies associated with big data also highlight such hurdles as data integration, data quality, data confidentiality and complexity of analysis procedure. [5] The present paper is different because it is directed toward a practical solution comprising all the stages of the pipeline for the analysis of structured financial statement datasets.

VI. PROPOSED SYSTEM

The system under consideration is divided into stages through which the user will be able to upload a financial dataset having a variety of fields like period, statement type, metric, and value. After that, the extraction component is supposed to go through this data and identify valid financial records. After extraction, the cleaning component would process the data through missing values and duplicate records, as well as deal with inconsistent labels and numbers formatting. Afterwards, it will be possible to apply standardization, whereby similar financial metrics will be assigned the same name.

VII. SYSTEM ARCHITECTURE

Stages	Elementary	Function
Processing	To upload the financial statements information.	
Extraction	To perform scanning and to recognize financial documents.	
Cleaning	To clean the data from emptiness, duplication, and unverifiable data.	
Standardization	To match against the given names, timeframes, and amounts.	
Saving	To save clear data.	



VII. METHODOLOGY

7.1 Uploading and extracting of data. The system accepts financial data that is structured and contains information about the reporting period, type of statement, kind of financial metric in figures. The extraction process checks against expected fields of the constraints and transforms incoming data to the internal format. 7.2 Data cleaning. The cleaning phase must be completed before data processing so that invalid or inconsistent values do not impact the calculations. Major actions usually performed during this phase are elimination of duplicates, processing of empty values, conversion of numeric parameters, and validation of the types of statements. 7.3 Standardization. The financial data may differ by the names in various sets of data. In this solution, the system maps these names to standard ones. Reporting periods are also standardized so that they can be easily compared chronologically. 7.4 Storing. Clean records will be stored in a


structured database. Separation of storage and input files will give an opportunity for analysis without repeated operations. 7.5 Ratio analysis. The system calculates financial ratios from the standardized data in the output format. Some examples may include comparisons of current ratios, debt equity ratio, net profit margins, and return on assets..

VIII.ALGORITHM





The system also provides what-if scenario simulation by allowing users to modify financial parameters such as revenue, expenses, and debt and recalculating the resulting financial metrics. Explainable AI techniques are then used to identify and explain the factors contributing to financial risks, anomalies, and changes in performance. Finally, the system generates meaningful financial insights and presents KPIs, ratios, trends, risks, anomalies, scenario comparisons, and AI-generated explanations through an interactive dashboard, with the option to generate and download a comprehensive financial analysis report. The system also provides what-if scenario simulation by allowing users to modify financial parameters such as revenue, expenses, and debt

and recalculating the resulting financial metrics. Explainable AI techniques are then used to identify and explain the factors contributing to financial risks, anomalies, and changes in performance. Finally, the system generates meaningful financial insights and presents KPIs, ratios, trends, risks, anomalies, scenario comparisons, and AI-generated explanations through an interactive dashboard, with the option to generate and download a comprehensive financial analysis report. The system also offers a what-if scenario simulation feature to enable users to alter financial variables like revenue, expenses, and debt and obtain recalculations of the financial metrics. Afterwards, Explainable AI methods are employed for revealing and clarifying the elements leading to financial risks, irregularities, and deviations in performance. Finally, the system provides valuable financial insights and allows the presentation of key performance indicators, ratios, tendencies, risks, irregularities, scenario comparisons, and AI-generated explanations on an interactive board along with generating and downloading a

lengthy report on financial analysis.




before any ratio is calculated. This separation is important because financial analysis is only as reliable as the underlying data. The pipeline therefore treats validation and standardization as first-class stages rather than optional preprocessing.



IX. IMPLEMENTATION



The project can be implemented using a web-based application architecture. A frontend provides the upload interface and analytical dashboard, while a backend handles validation, transformation, database operations and ratio calculations. A relational database can be used to store standardized financial records. The implementation is intentionally modular so that additional statement types, ratios and visualization components can be added without redesigning the entire pipeline. The dataprocessing layer is the core component. It transforms uploaded records into a consistent schema



X. RESULTS AND DISCUSSION

The developed system shows that an orderly and organized algorithm for data processing can simplify the process of analyzing financial statements. If financial records are made uniform, ratios can be calculated consistently for different periods. The trends obtained will help in better understanding the changes in liquidity, profitability, solvency, and operational efficiency during the periods under review. The main contribution of the project is not a new financial ratio but an integration of several analytical operations into a single process. Classic ratio analysis is useful; however, the research points out that ratios provide only a part of the information contained in financial statements. For this reason, the mentioned system aims at presenting ratios together with other relevant financial indicators as well as trends in various periods. For a student

project, this pipeline will also demonstrate a practical implementation of the following data science skills: data ingestion, data validation, ETL, database management, backend calculations, API-based communication, and data visualization.



XII. LIMITATIONS

The correctness of the analysis results relies on how good and reliable the financial data is. The initial version mainly works with structured datasets and does not work especially well with complex PDFs, handwritten texts, or incredibly unstructured financial disclosures. The reading of financial ratios also needs to be performed with caution since

it is possible that accounting practice, industry specifics, and size of the company are capable of influencing their meanings.

XIV.CONCLUSION

The article provides a financial statement analytics pipeline that connects the process of financial data extraction, transformation, storage, and analysis. The system serves as a good example of using software engineering together with financial analysis to automate the processes and produce the same analytical results with minimum efforts. With the help of this tool, it will be easy to analyze financial ratios and compare tendencies across periods. Several options for extending the system include the incorporation of intelligent financial analytics tools, such as document extraction, predictive modeling and anomaly detection. In the future, the integration of automated document extraction from PDF reports with the use of OCR and document recognition technologies will make the process even easier. Additionally, anomaly detection technologies will be implemented for the identification of outliers and prediction of financial risks for the execution of automated financial forecasting and

monitoring. Automated benchmarking tools will also be implemented to compare companies with each other.

XV. LITERATURE SURVEY

1. Financial Big Data Analysis and Early Warning Platform: A Case Study – Liang et al. (2020) The paper presents a big-data platform for financial risk analysis and early warning. It combines heterogeneous data collection, data cleaning, data fusion, cloud computing, and AI to improve financial risk monitoring and decision-making.

2. Intelligent Fraud Detection in Financial Statements Using Machine Learning and Data Mining – Ashtiani & Raahemi (2021) This paper reviews machine learning and data-mining techniques used for detecting fraudulent financial statements. The study analyzes 47 research papers and highlights the importance of anomaly detection, unstructured financial data, and advanced ML techniques for future fraud detection systems.

3. Financial Knowledge Graph Based Financial Report Query System – Zehra et al. (2021) The authors developed a financial knowledge graph using

ontologybased information extraction to extract useful information from annual financial reports. The system enables users to query financial information and automatically generate financial stories for decision-making.

4. Sentiment Analysis in Finance: From Transformers Back to eXplainable Lexicons (XLex) – Rizinski et al. (2024) This paper proposes XLex, which combines transformer models and SHAP explainability to improve financial sentiment analysis. The method improves financial vocabulary coverage and classification accuracy while being more efficient and interpretable than conventional transformer models.

5. AI Driven Fraud Detection Models in Financial Networks: A Comprehensive Systematic Review – Sarna et al. (2025) This systematic review examines AI, machine learning, deep learning, hybrid models, and GNNs for financial fraud detection. It highlights the benefits of AI for identifying fraud patterns while identifying explainability, privacy, bias, and scalability as major challenges.

6. Multimodal Fraud Detection in Financial Statements: A Trimodal

Attention Network With Contrastive Evidence Chain Construction – Ping et al. (2026)

The paper proposes TriAudit, which combines financial ratios, textual information, and document-layout features for financial statement fraud detection. The multimodal approach improves fraud detection performance while providing interpretable evidence linking financial, textual, and visual anomalies. .

7. Hajek, P. and Henriques, R. (2017), “Mining Corporate Annual Reports for Intelligent Detection of Financial Statement Fraud – A Comparative Study of Machine Learning Methods.”

Main Contribution: Combined financial information and managerial comments with machine-learning methods for financial-statement fraud detection. Relevance: Highly relevant to the proposed integration of structured financial data and unstructured textual information.

8. “A Research on Text Information Extraction from Annual Report Based on Domain Ontology” (2017). Main Contribution: Applied domain ontology and naturallanguage-

processing techniques to extract information from annual reports. Relevance: Supports the automated financial-document extraction and data standardization stages of the proposed pipeline.

9. Ito, T., Sakaji, H. and Izumi, K. (2019), "Segment Information Extraction from Financial Annual Reports Using Neural Network." Main Contribution: Applied neural-network techniques to extract segment information from financial annual reports. Relevance: Supports automated extraction of structured financial information from unstructured annual reports.

10. Araci, D. (2019), "FinBERT: Financial Sentiment Analysis with Pre-trained Language Models." Main Contribution: Introduced a BERT-based language model specialized for financial-domain text. Relevance: Provides a foundation for applying domain-specific NLP to financial documents and generating meaningful financial insights.

11. Chen, Z. et al. (2021), "FinQA: A Dataset of Numerical Reasoning over Financial Data." Main Contribution: Introduced a financial question-answering

dataset requiring numerical reasoning over financial reports. Relevance: Supports the numerical reasoning and AI-based financial insight generation components of the proposed system.

12. Zhu, F. et al. (2021), "TATQA: A Question Answering Benchmark on a Hybrid of Tabular and Textual Content in Finance." Main Contribution: Developed a benchmark combining tabular and textual financial information and requiring numerical reasoning. Relevance: Supports the integration of extracted financial tables and narrative information within the proposed analytics pipeline.

13. "On the Determinants and Prediction of Corporate Financial Distress in India" (2021). Main Contribution: Investigated financial ratios and machine learning models such as SVM and ANN for corporate financial distress prediction in India. Relevance: Provides evidence for using financial indicators and machine learning for financial risk prediction in the Indian context.

14. "Predicting Financial Distress of Chinese Listed Companies

Using Machine Learning: To What Extent Does Textual Disclosure Matter?” (2023). Main Contribution: Investigated whether textual disclosures combined with financial information improve financial distress prediction. Relevance: Demonstrates the importance of combining structured financial metrics with unstructured textual disclosures.

15. “Assessing Large Language Models Used for Extracting Table Information from Annual Financial Reports” (2024). Main Contribution: Investigated large language models for extracting and interpreting tabular information from annual financial reports. Relevance: Directly supports the use of modern AI/LLM approaches for automated

Using Machine Learning: To What Extent Does Textual Disclosure Matter?” (2023). Main Contribution: Investigated whether textual disclosures combined with financial information improve financial distress prediction. Relevance: Demonstrates the importance of combining structured financial metrics with unstructured textual disclosures.

15. “Assessing Large Language Models Used for Extracting Table Information from Annual Financial Reports” (2024). Main Contribution: Investigated large language models for extracting and interpreting tabular information from annual financial reports. Relevance: Directly supports the use of modern AI/LLM approaches for automated

Match Groups

- 51 Not Cited or Quoted 11%**
Matches with neither in-text citation nor quotation marks
- 2 Missing Quotations 1%**
Matches that are still very similar to source material
- 0 Missing Citation 0%**
Matches that have quotation marks, but no in-text citation
- 0 Cited and Quoted 0%**
Matches with in-text citation present, but no quotation marks

Top Sources

- 5% Internet sources
- 7% Publications
- 0% Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1	Publication	Financial fraud detection through the application of	2%
2	Internet	urfpublishers.com	1%
3	Publication	using OCAI. Frontiers in Psychology, 14, p.1297041	1%
4	Publication	Office of Maternal and Child Health Active Projects FY 1989. An Annotated Li ...	1%
5	Publication	THE DETERMINANTS OF BEHAVIOURAL INTENTION TO ADOPT E-WALLET AMONG THE ...	1%
6	Publication	A Systematic Review of Cyber Threat Intelligence: The Effectiveness of Techn ...	1%
7	Publication	An Analysis on Financial Statement Fraud Detection for Chinese Listed Compan ...	1%
8	Publication	Large Language Models for Financial and Investment Management: Applications ...	1%
9	Publication	IMPACT OF AUDITING ON THE INTEGRITY OF FINANCIAL STATEMENTS OF QUOTED CO ...	1%
10	Internet	naturelightpublications.com	1%

11	Internet	www.semantic-web-journal.net	1%
12	Internet	projekter.aau.dk	<1%
13	Internet	global.asrcconference.com	<1%
14	Internet	wsr-j.org	<1%
15	Internet	www.fathomhq.com	<1%
16	Internet	www.avepubs.com	<1%
17	Publication	Intelligent Financial Statement Extraction from Complex PDF Documents Using ...	<1%
18	Publication	Youth Development: A Case Study from Honduras.	<1%
19	Publication	A Survey of Explainable NLP Across Various Domains	<1%
20	Publication	Explainable machine learning framework for assessing groundwater quality and ...	<1%
21	Internet	iris.unil.ch	<1%
22	Internet	www.tqmp.org	<1%
23	Internet	optimization-online.org	<1%
24	Internet	aclanthology.org	<1%
25	Internet	accountants.sva.com	<1%

26	Publication	Graph neural networks for financial fraud detection: a review	<1%
27	Internet	ojs.ukscip.com	<1%
28	Publication	Artificial Intelligence Driven Approaches for Financial Fraud Detection: A S ...	<1%
29	Internet	ijcope.org	<1%
30	Publication	Perceived Influence of Teachers' Guidance on Student's Choice of Advanced Le ...	<1%
31	Publication	Bridging language models and financial analysis	<1%
32	Publication	Next-Generation Fraud Detection System Using AI And Blockchain for Real-Time ...	<1%
33	Publication	[8] Yang, H., Liu, X.-Y., & Wang, C. D. (2023). FinGPT: Open-source financia ...	<1%
34	Publication	Corporate financial distress prediction with multiperiod annual report data: ...	<1%
35	Internet	www.ukm.my	<1%
36	Internet	ou.ac.lk	<1%
37	Publication	Luna: an evaluation foundation model to catch language model hallucinations ...	<1%
38	Publication	Backchaining Loss of Control Mitigations from Mission-Specific Benchmarks in ...	<1%